

# Introduction to Unsupervised Learning

# Types of Machine Learning

Supervised

data points have known outcome

Unsupervised

data points have unknown outcome

# Types of Machine Learning

Supervised

data points have known outcome

Unsupervised

data points have unknown outcome

# Types of Unsupervised Learning

Clustering

identify unknown structure in data

# Types of Unsupervised Learning

Clustering

identify unknown structure in data

Dimensionality  
Reduction

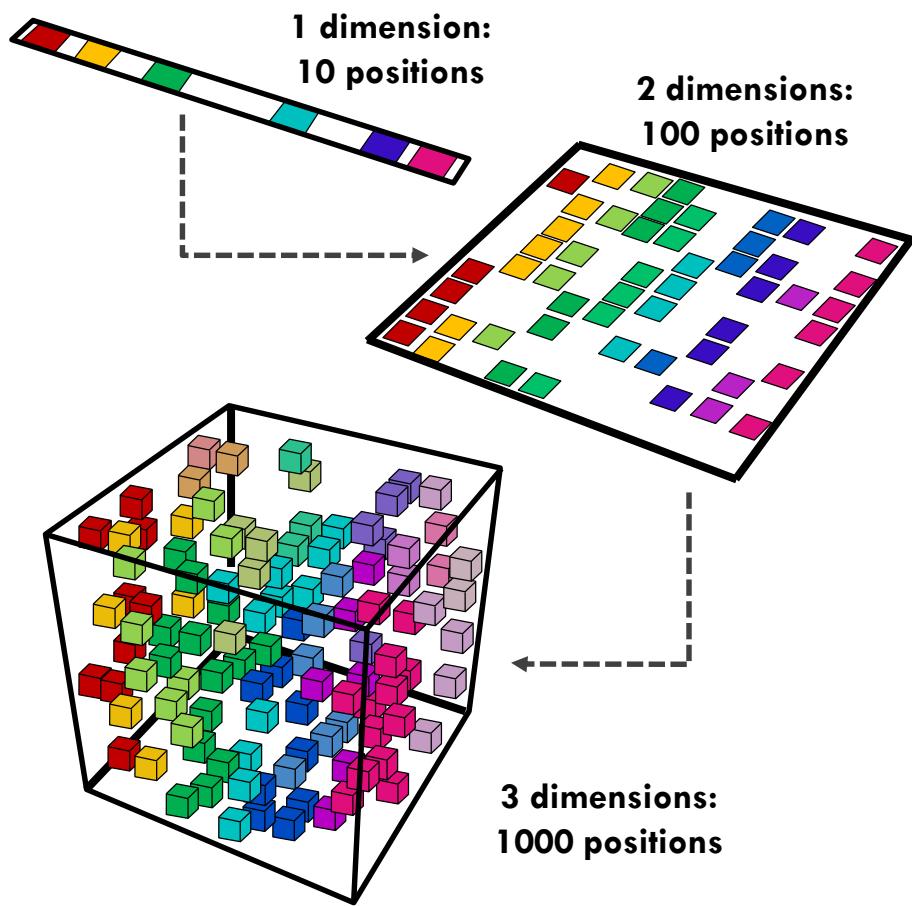
use structural characteristics to simplify data

---

# Dimensionality Reduction

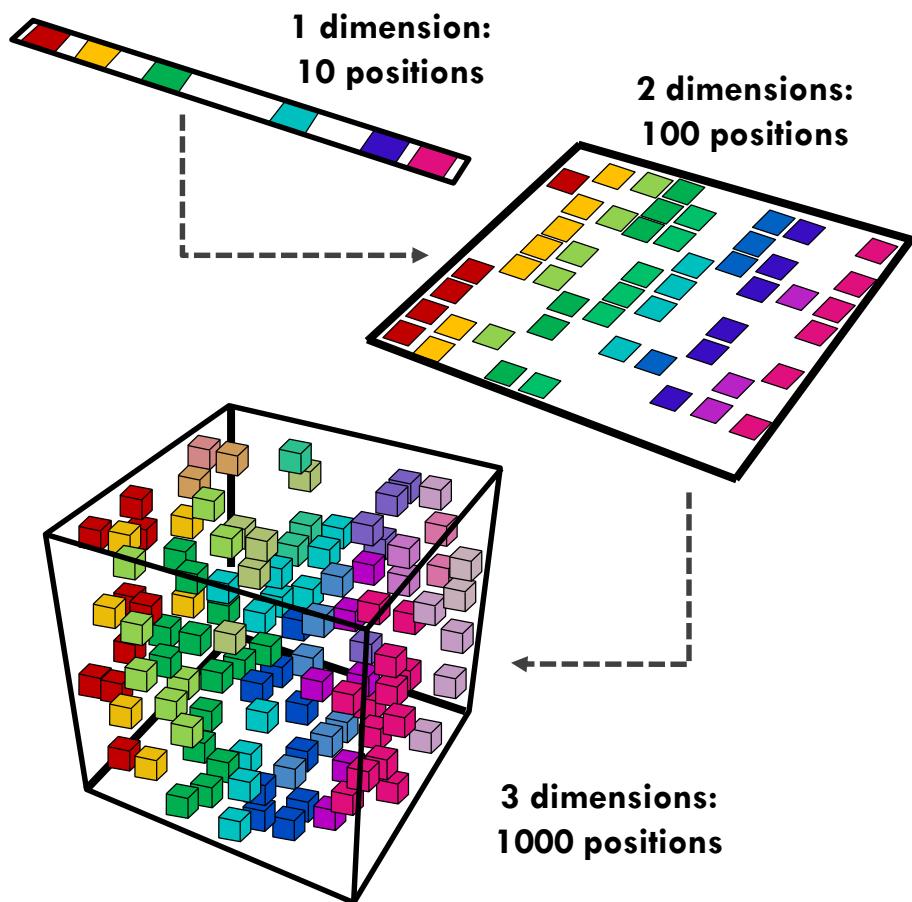
# Curse of Dimensionality

- Theoretically, increasing features should improve performance



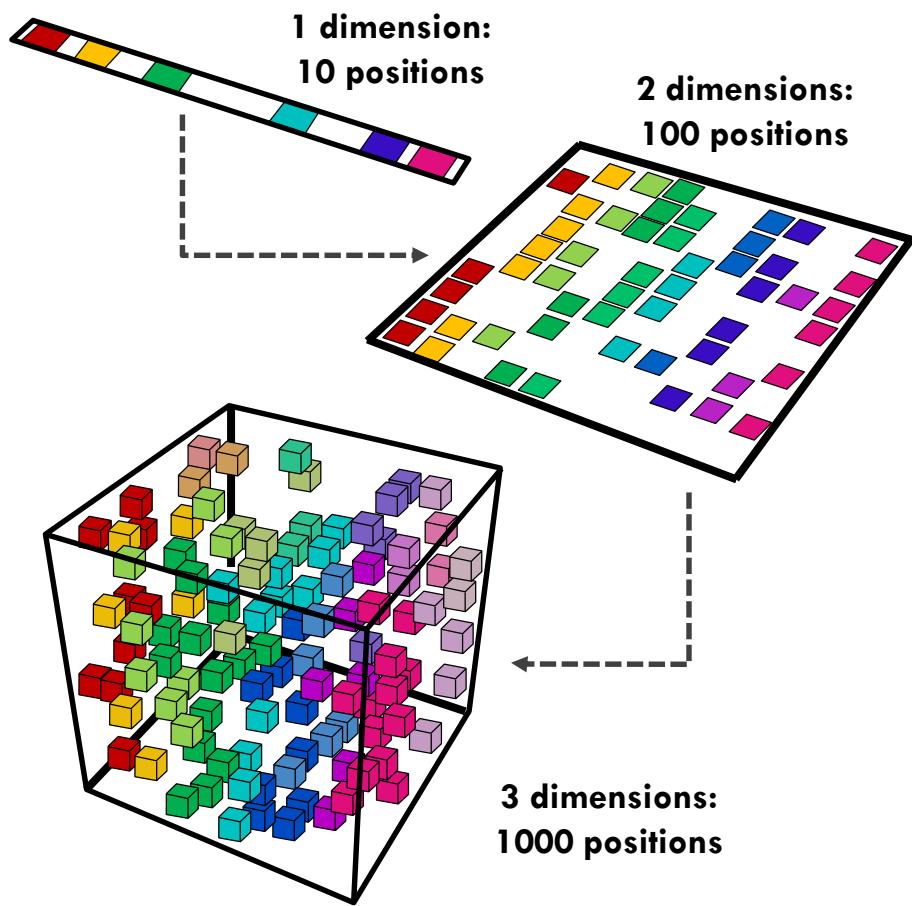
# Curse of Dimensionality

- Theoretically, increasing features should improve performance
- In practice, too many features leads to worse performance



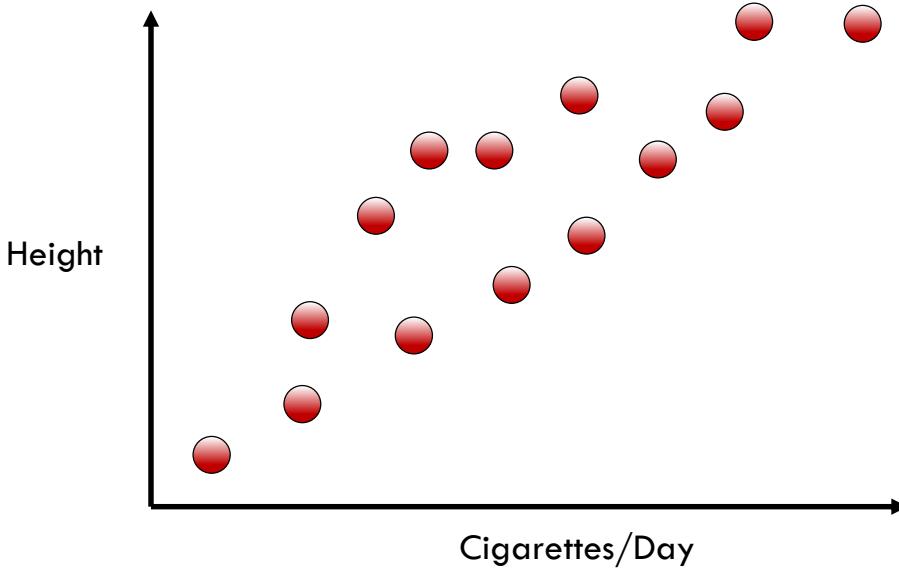
# Curse of Dimensionality

- Theoretically, increasing features should improve performance
- In practice, too many features leads to worse performance
- Number of training examples required increases exponentially with dimensionality



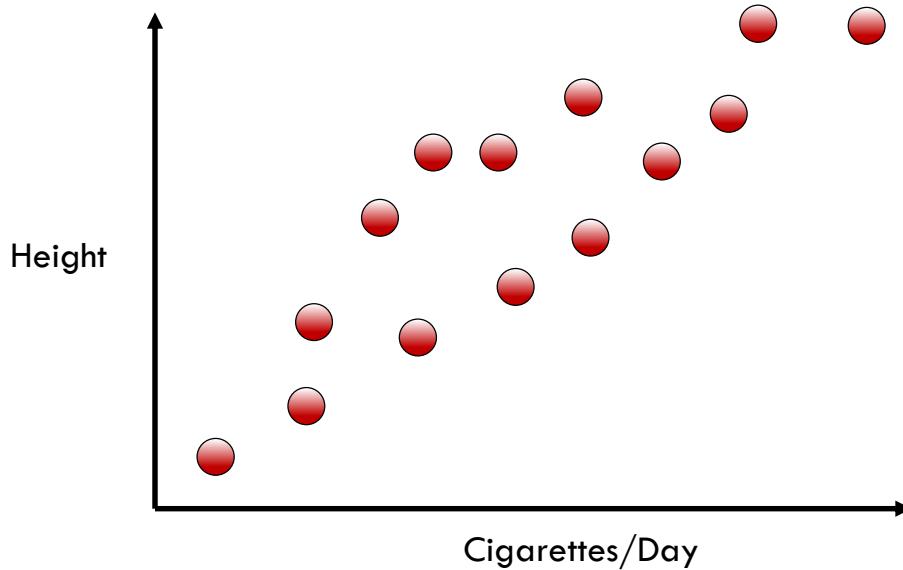
# Solution: Dimensionality Reduction

- Data can be represented by fewer dimensions (features)
- Reduce dimensionality by



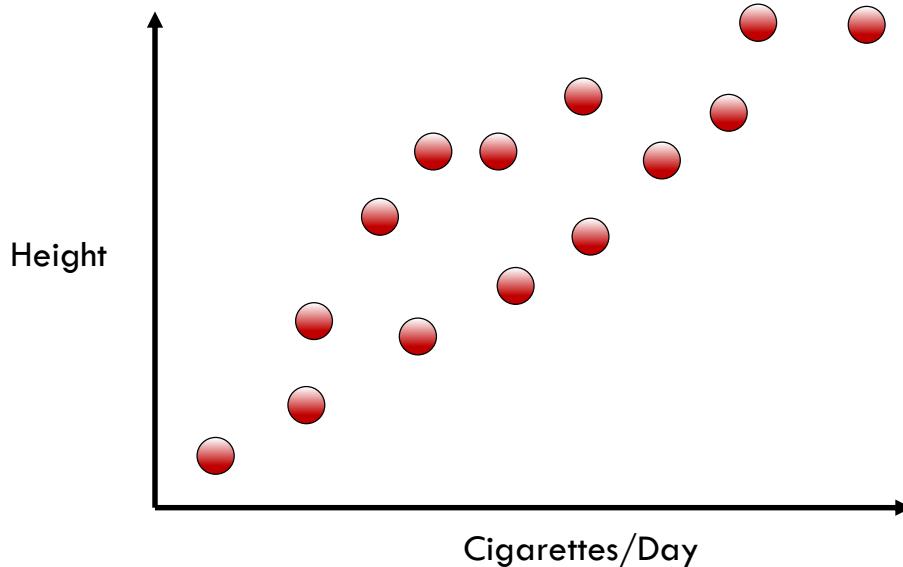
# Solution: Dimensionality Reduction

- Data can be represented by fewer dimensions (features)
- Reduce dimensionality by selecting subset (feature elimination)
- Combine with linear and non-



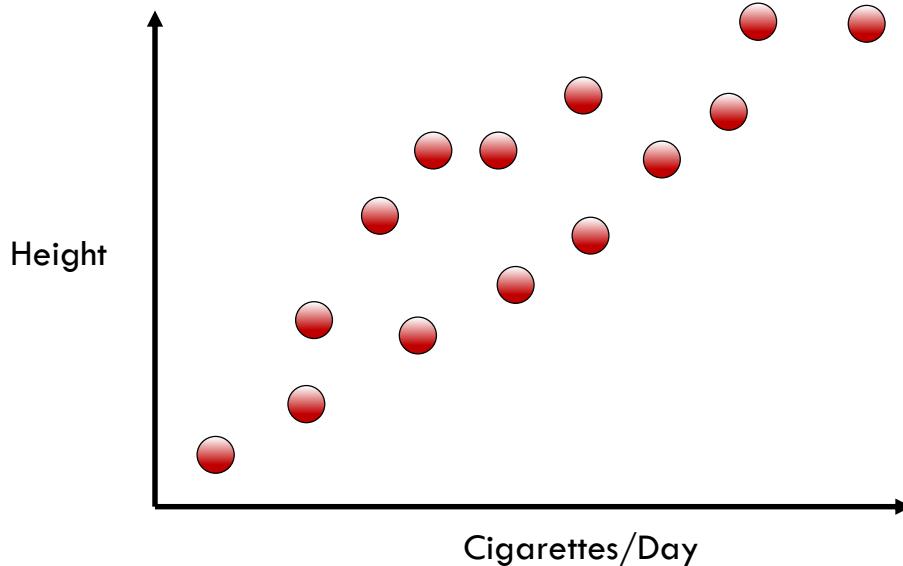
# Solution: Dimensionality Reduction

- Data can be represented by fewer dimensions (features)
- Reduce dimensionality by selecting subset (feature elimination)
- Combine with linear and non-linear transformations



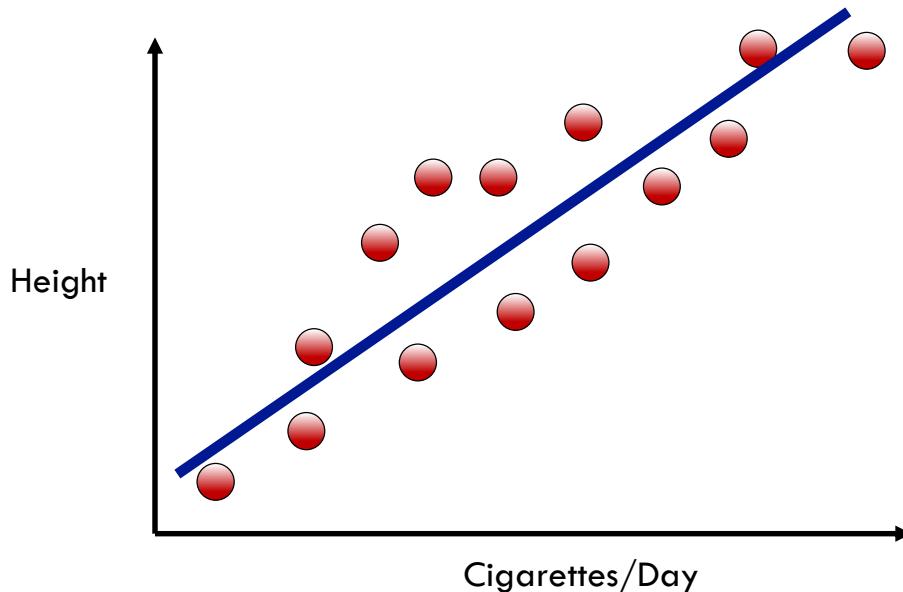
# Solution: Dimensionality Reduction

- Two features: height and cigarettes per day
- Both features increase together (correlated)
- Can we reduce number of features to one?



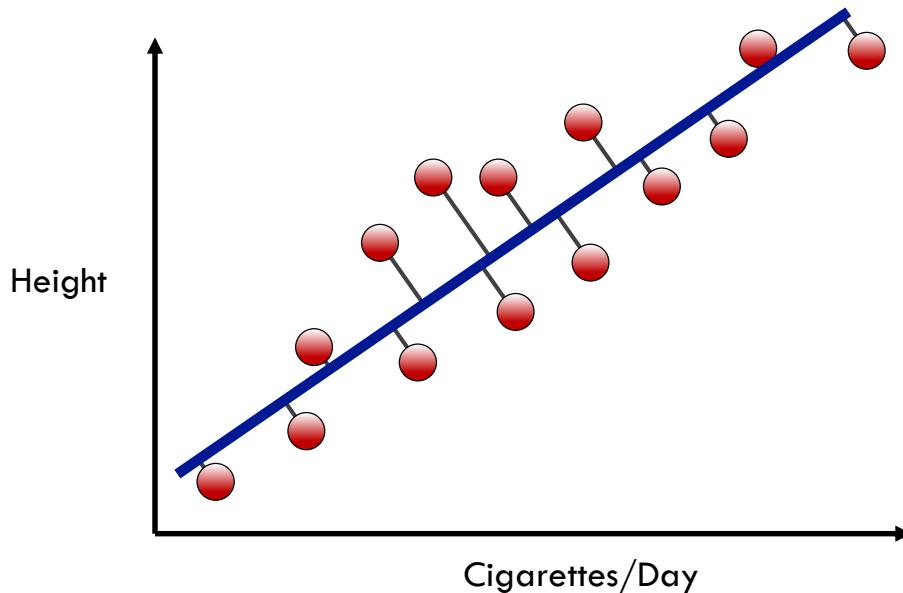
# Solution: Dimensionality Reduction

- Two features: height and cigarettes per day
- Both features increase together (correlated)
- Can we reduce number of features to one?



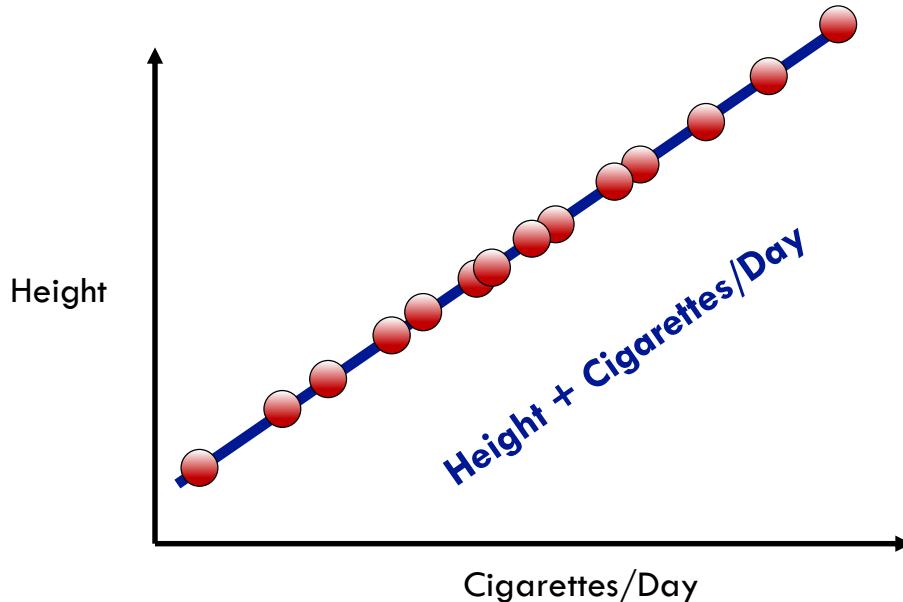
# Solution: Dimensionality Reduction

- Two features: height and cigarettes per day
- Both features increase together (correlated)
- Can we reduce number of features to one?



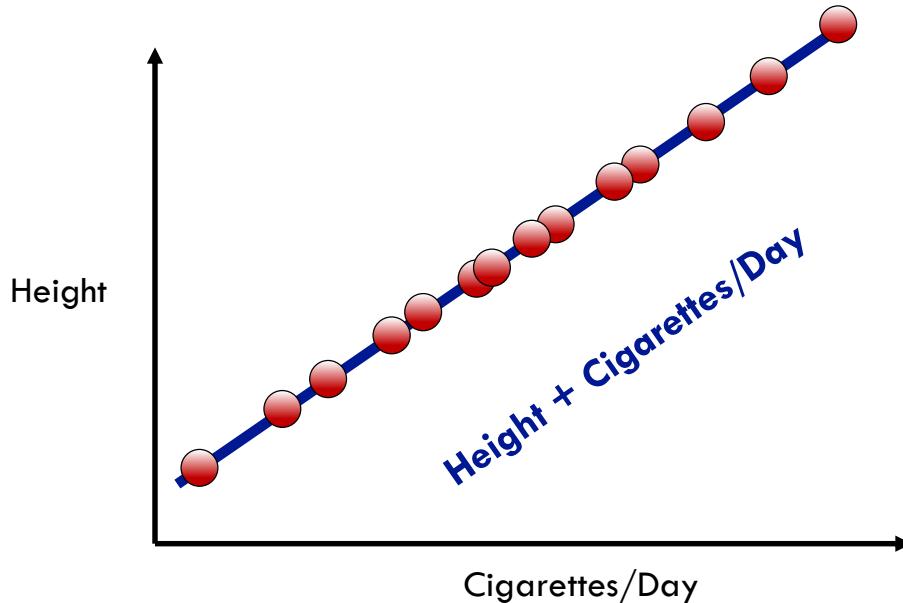
# Solution: Dimensionality Reduction

- Two features: height and cigarettes per day
- Both features increase together (correlated)
- Can we reduce number of features to one?



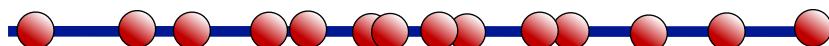
# Solution: Dimensionality Reduction

- Create single feature that is combination of height and cigarettes
- This is Principal Component Analysis (PCA)



# Solution: Dimensionality Reduction

- Create single feature that is combination of height and cigarettes
- This is Principal Component Analysis (PCA)



Height + Cigarettes/Day

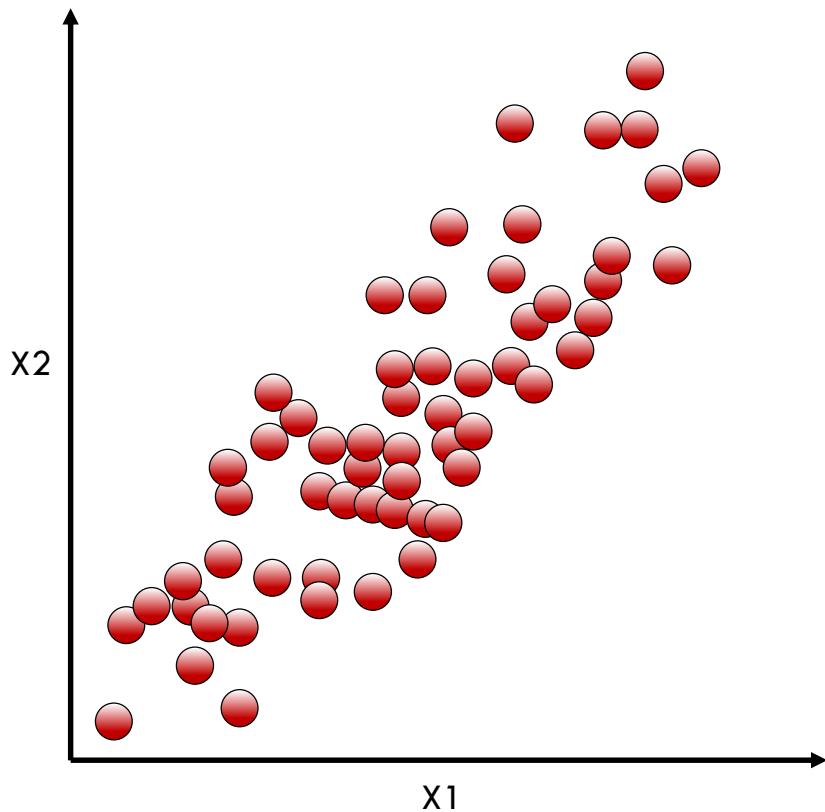
# Dimensionality Reduction

Given an  $N$ -dimensional data set ( $x$ ), find a  $N \times K$  matrix ( $U$ ):

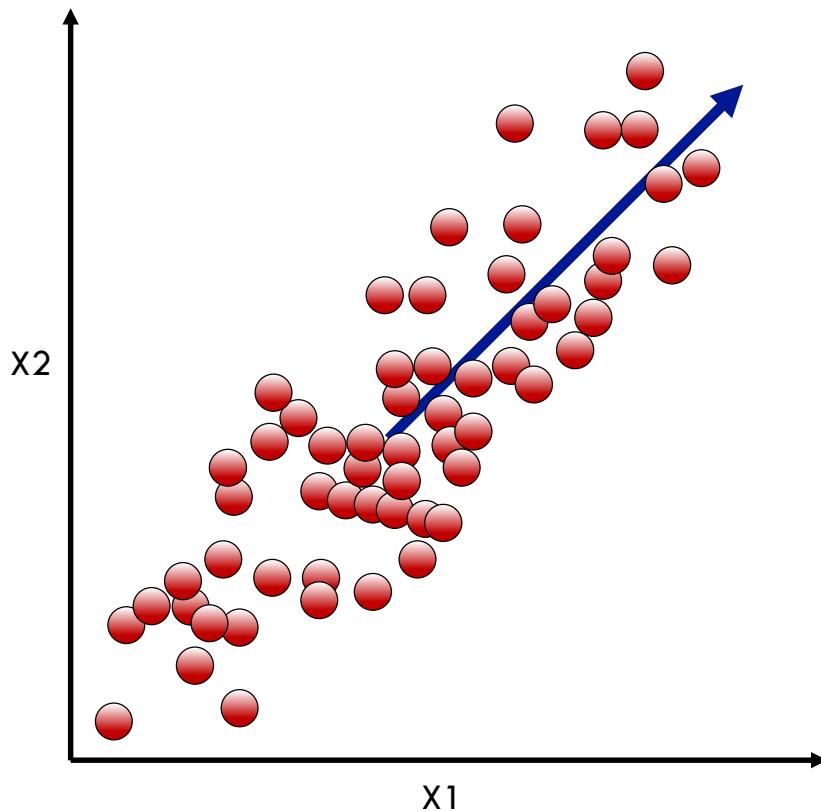
$y = U^T x$ , where  $y$  has  $K$  dimensions and  $K < N$

$$x = \begin{bmatrix} x_1 \\ x_2 \\ \dots \\ x_n \end{bmatrix} \xrightarrow{U^T} y = \begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_k \end{bmatrix} (K < N)$$

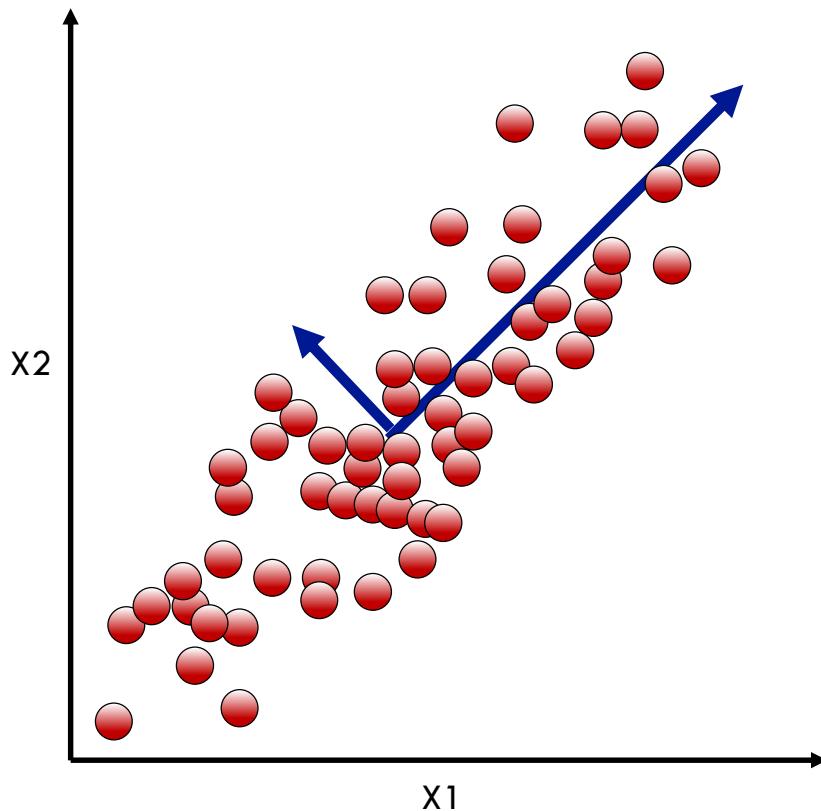
# Principal Component Analysis (PCA)



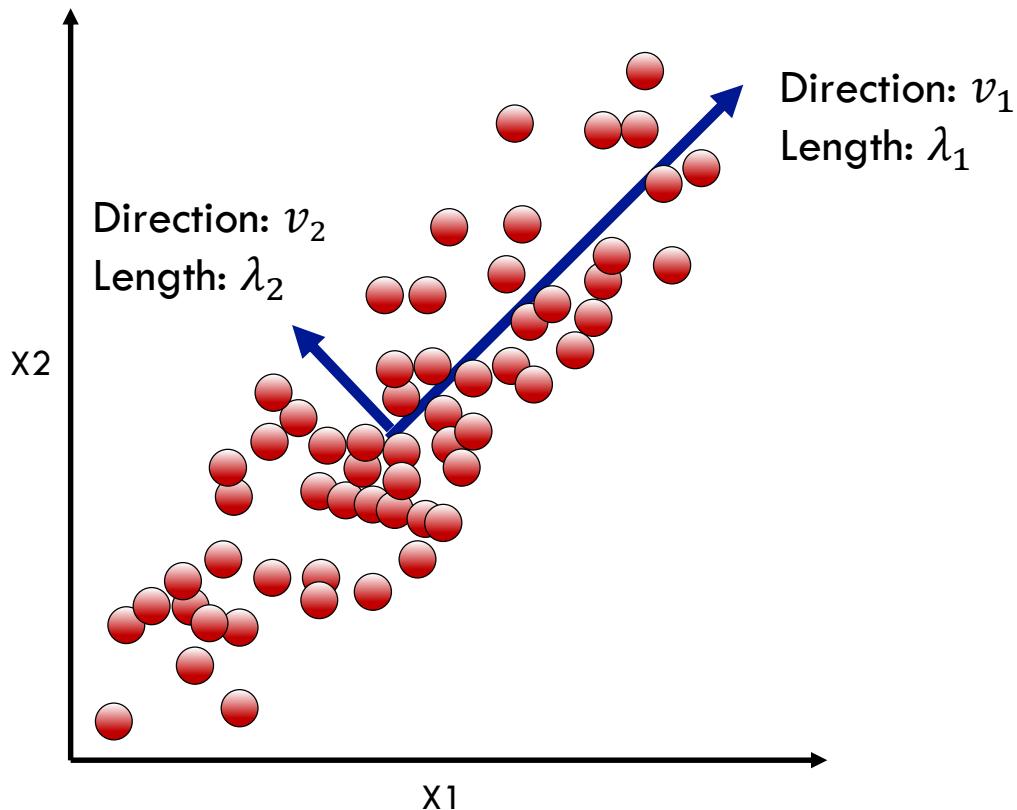
# Principal Component Analysis (PCA)



# Principal Component Analysis (PCA)



# Principal Component Analysis (PCA)



# Single Value Decomposition (SVD)

- SVD is a matrix factorization method normally used for PCA
- Does not require a square data set
- SVD is used by Scikit-learn for PCA

$$\begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \\ * & * & * \\ * & * & * \end{bmatrix} = \begin{bmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{bmatrix} \begin{bmatrix} * & 0 & 0 \\ 0 & * & 0 \\ 0 & 0 & * \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \end{bmatrix}$$
$$A_{m \times n} \quad U_{m \times m} \quad S_{m \times n} \quad V_{n \times n}^T$$

# Truncated Single Value Decomposition

- How can SVD be used for dimensionality reduction?
- Principal components are calculated from  $US$
- "Truncated SVD" used for dimensionality reduction ( $n \rightarrow k$ )

$$A_{m \times n} \approx U_{m \times k} \begin{bmatrix} 9 & 0 & 0 \\ 0 & 7 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} V_{k \times n}^T$$

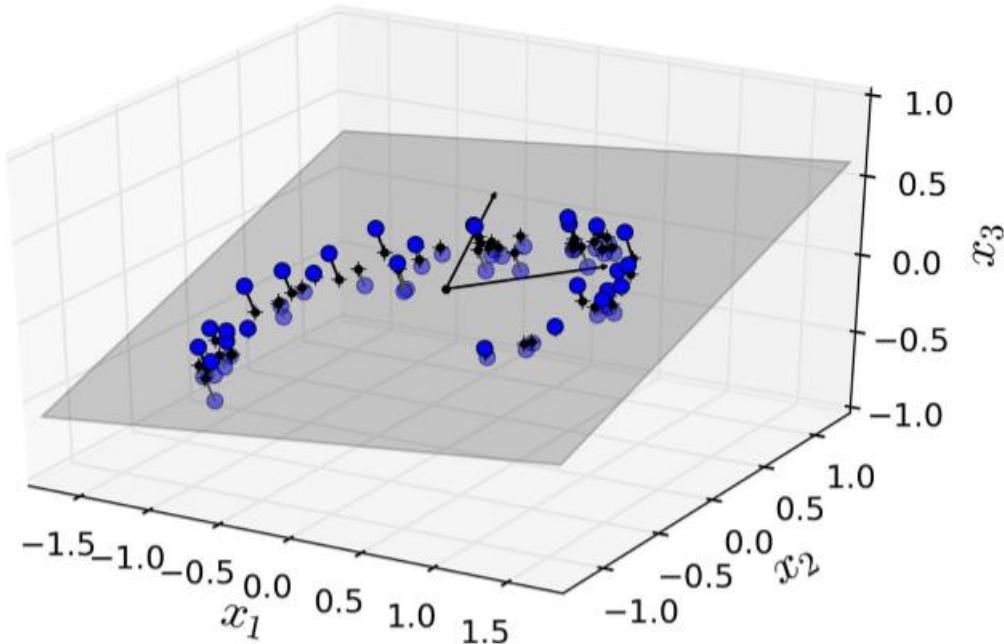


Figure 8-2. A 3D dataset lying close to a 2D subspace

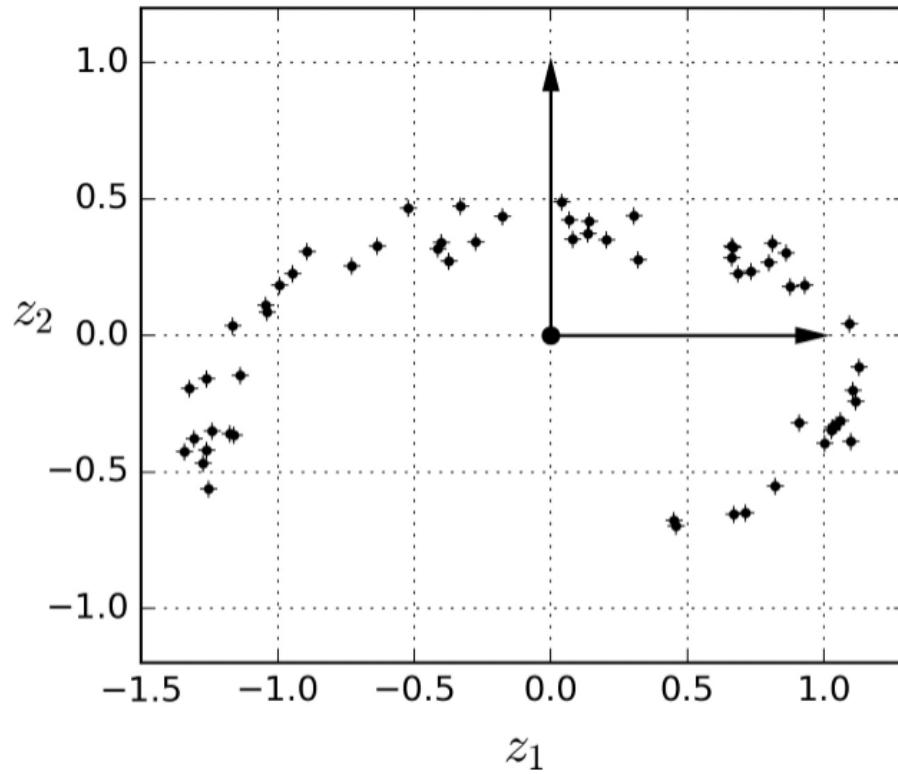
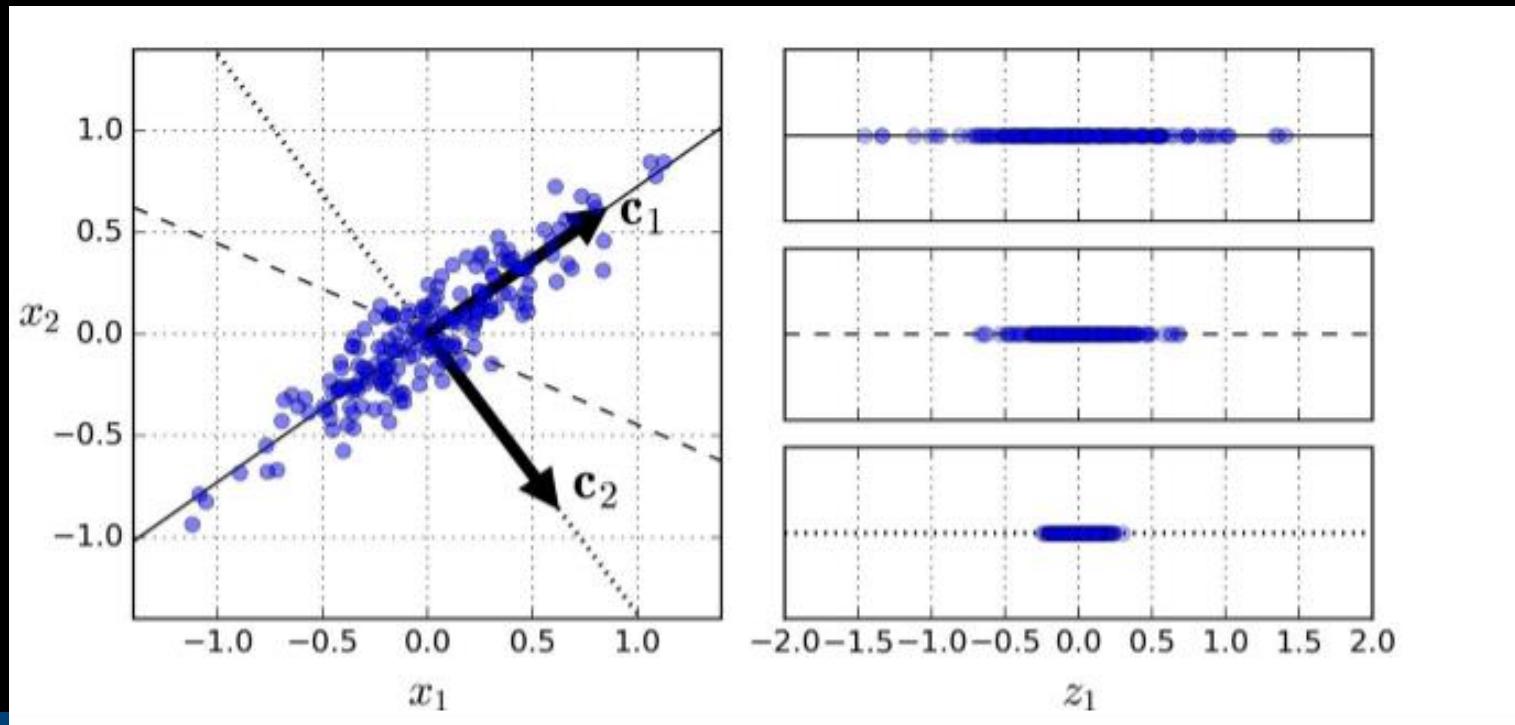


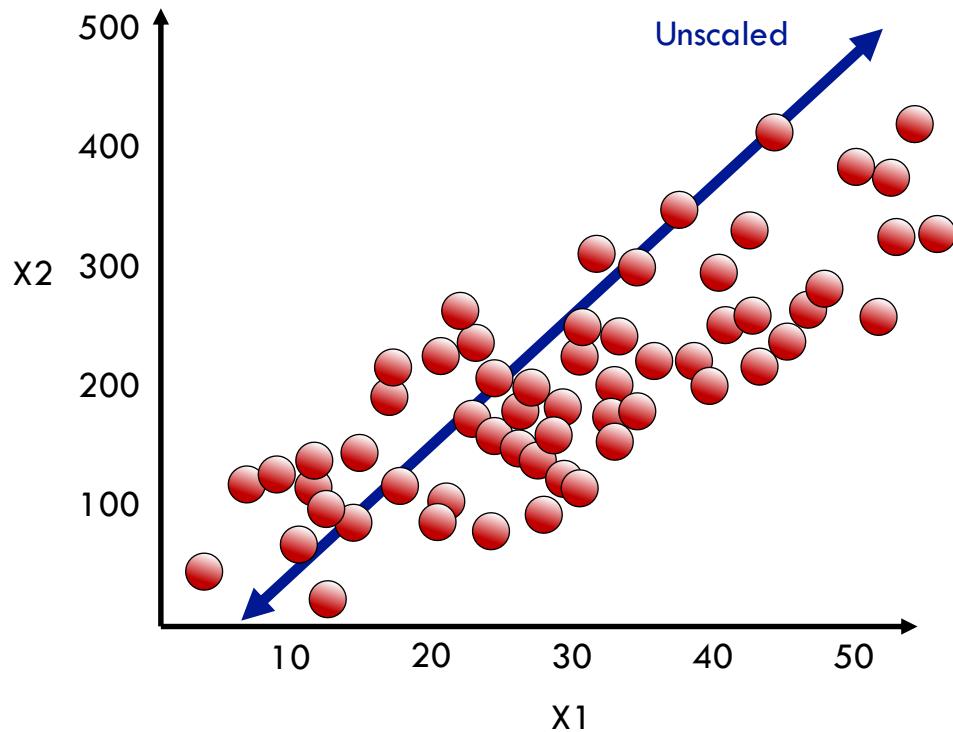
Figure 8-3. The new 2D dataset after projection

# Projections along various axis



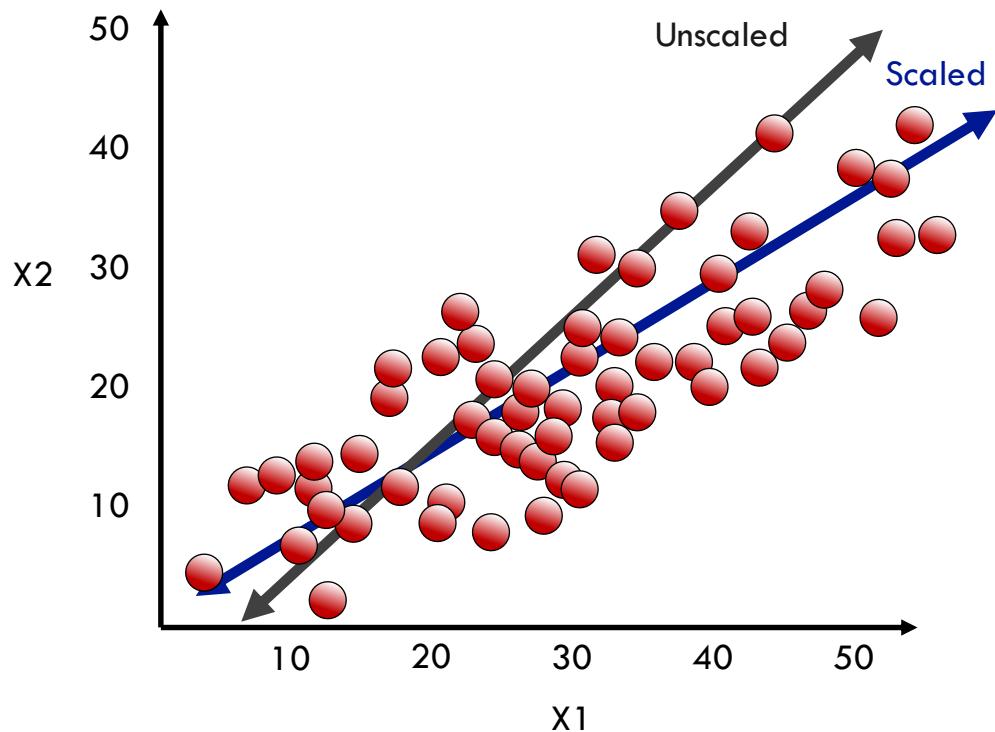
# Importance of Feature Scaling

- PCA and SVD seek to find the vectors that capture the most variance
- Variance is sensitive to axis scale



# Importance of Feature Scaling

- PCA and SVD seek to find the vectors that capture the most variance
- Variance is sensitive to axis scale
- Must scale data!



# PCA: The Syntax

**Import the class containing the dimensionality reduction method**

```
from sklearn.decomposition import PCA
```

# PCA: The Syntax

**Import the class containing the dimensionality reduction method**

```
from sklearn.decomposition import PCA
```

**Create an instance of the class**

```
PCAinst = PCA(n_components=3, whiten=True)
```

# PCA: The Syntax

**Import the class containing the dimensionality reduction method**

```
from sklearn.decomposition import PCA
```

**Create an instance of the class**

```
PCAinst = PCA(n_components=3, whiten=True)
```



final number of dimensions

# PCA: The Syntax

**Import the class containing the dimensionality reduction method**

```
from sklearn.decomposition import PCA
```

**Create an instance of the class**

```
PCAinst = PCA(n_components=3, whiten=True)
```



whiten = scale  
and center data

# PCA: The Syntax

**Import the class containing the dimensionality reduction method**

```
from sklearn.decomposition import PCA
```

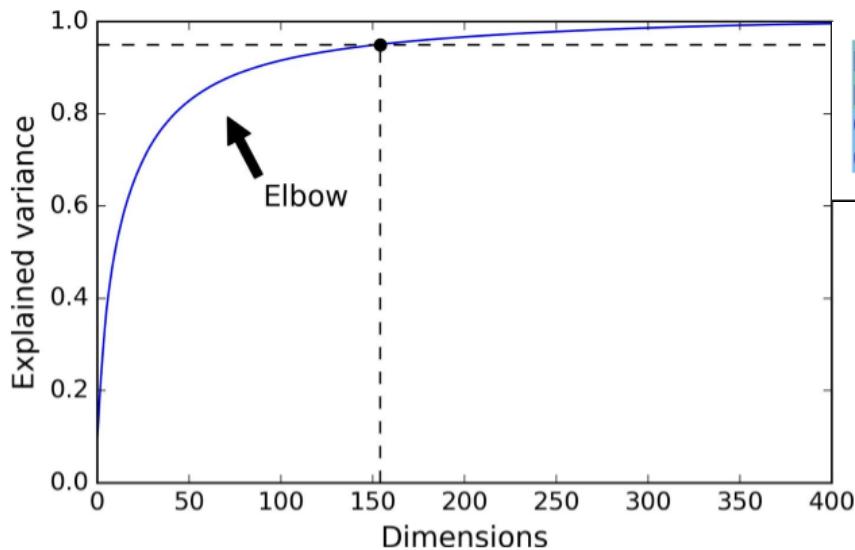
**Create an instance of the class**

```
PCAinst = PCA(n_components=3, whiten=True)
```

**Fit the instance on the data and then transform the data**

```
X_trans = PCAinst.fit_transform(X_train)
```

# Explained variance



```
pca = PCA()  
pca.fit(X)  
cumsum = np.cumsum(pca.explained_variance_ratio_)  
d = np.argmax(cumsum >= 0.95) + 1
```

# PCA: The Syntax

**Import the class containing the dimensionality reduction method**

```
from sklearn.decomposition import PCA
```

**Create an instance of the class**

```
PCAinst = PCA(n_components=3, whiten=True)
```

**Fit the instance on the data and then transform the data**

```
X_trans = PCAinst.fit_transform(X_train)
```

**Does not work with sparse matrices**

# Truncated SVD: The Syntax

**Import the class containing the dimensionality reduction method**

```
from sklearn.decomposition import TruncatedSVD
```

**Create an instance of the class**

```
SVD = TruncatedSVD(n_components=3)
```

**Fit the instance on the data and then transform the data**

```
X_trans = SVD.fit_transform(X_sparse)
```

**Works with sparse matrices—used with text data for Latent Semantic Analysis (LSA)**

# Truncated SVD: The Syntax

**Import the class containing the dimensionality reduction method**

```
from sklearn.decomposition import TruncatedSVD
```

**Create an instance of the class**

```
SVD = TruncatedSVD(n_components=3)
```



does not center data

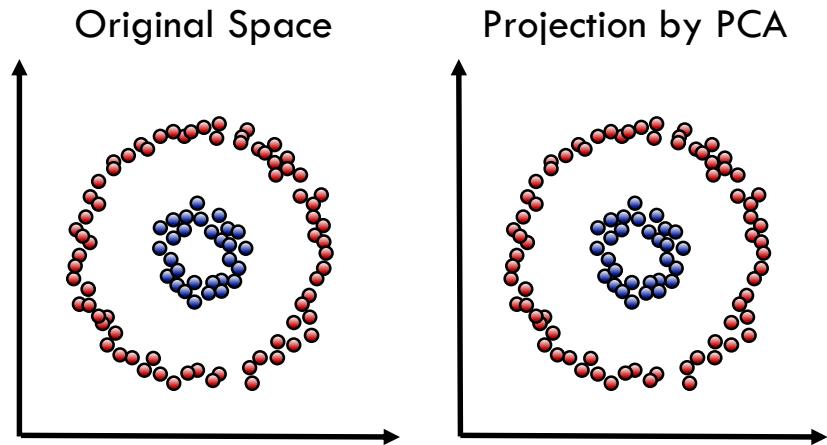
**Fit the instance on the data and then transform the data**

```
X_trans = SVD.fit_transform(X_sparse)
```

**Works with sparse matrices—used with text data for Latent Semantic Analysis (LSA)**

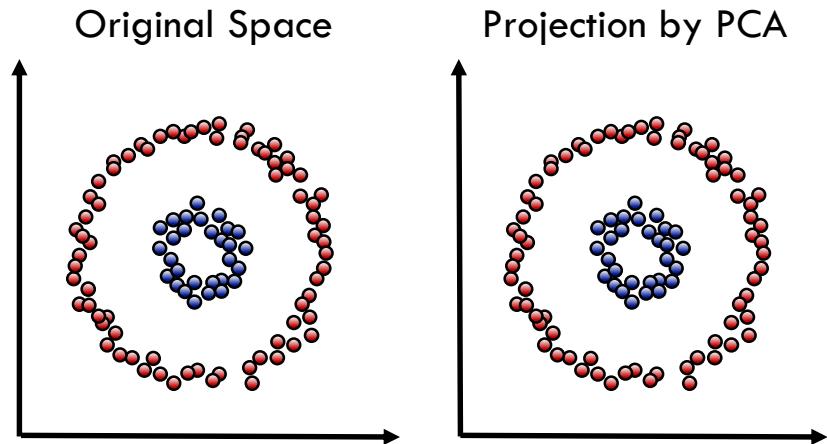
# Moving Beyond Linearity

- Transformations calculated with PCA/SVD are linear
- Data can have non-linear features
- This can cause dimensionality



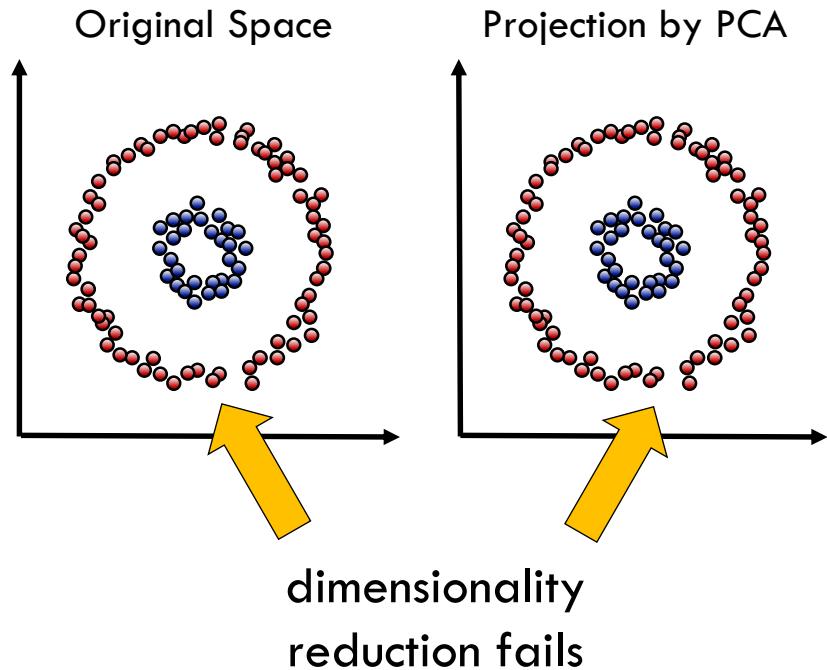
# Moving Beyond Linearity

- Transformations calculated with PCA/SVD are linear
- Data can have non-linear features
- This can cause dimensionality reduction to fail



# Moving Beyond Linearity

- Transformations calculated with PCA/SVD are linear
- Data can have non-linear features
- This can cause dimensionality reduction to fail



# Kernel PCA

- **Solution:** kernels can be used to perform non-linear PCA

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

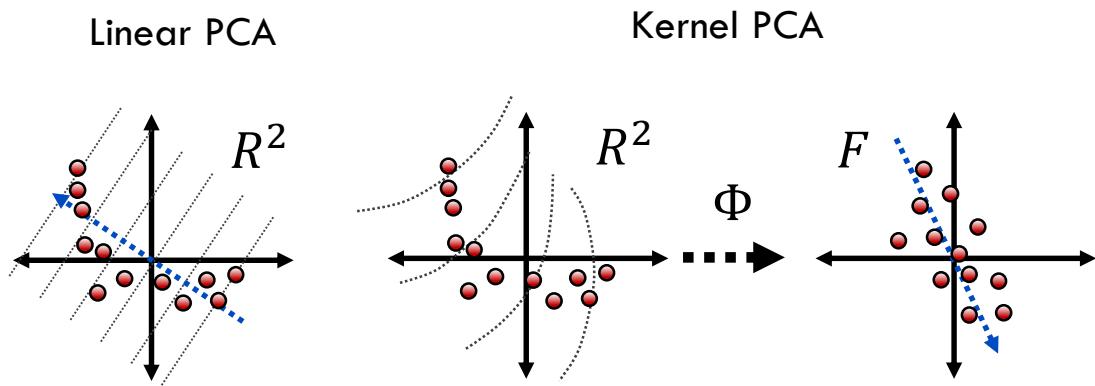
...

...

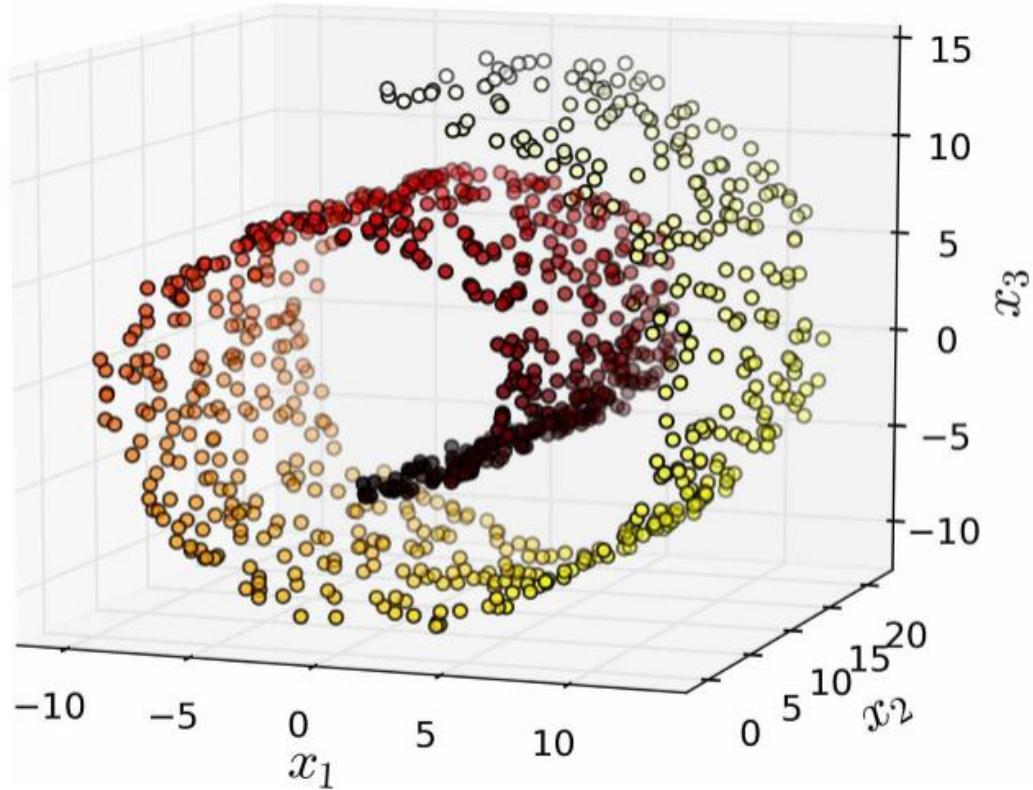
...

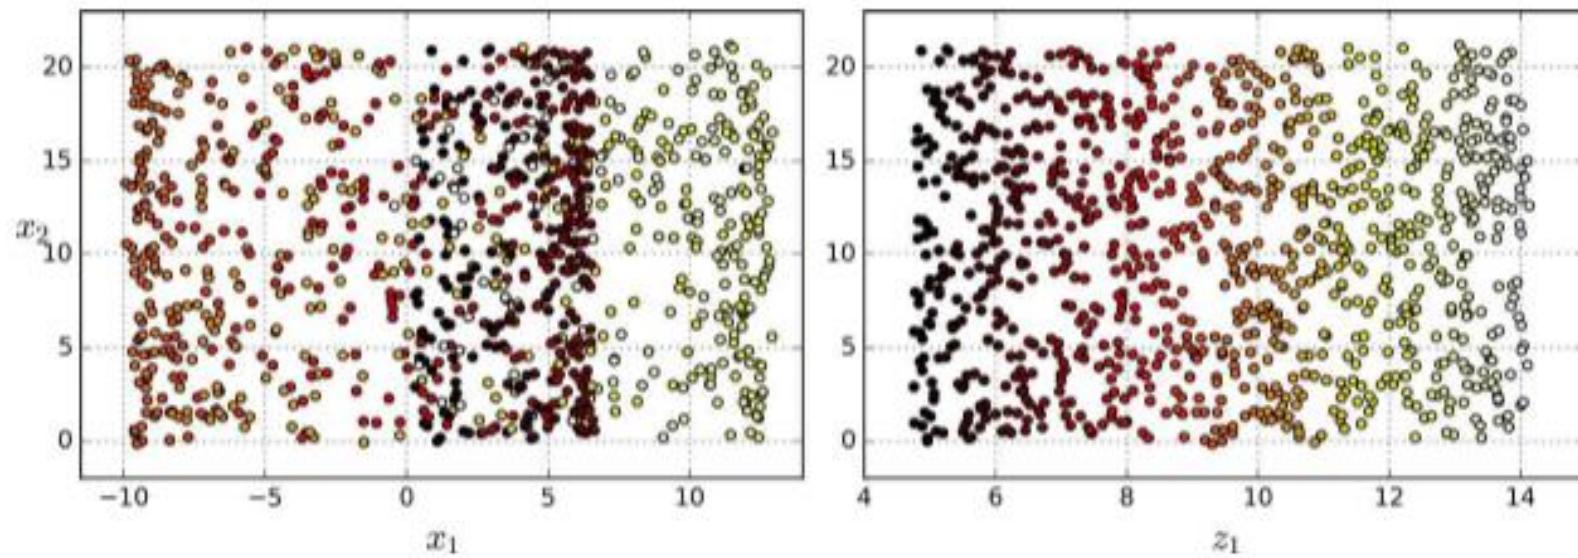
# Kernel PCA

- **Solution:** kernels can be used to perform non-linear PCA
- Like the kernel trick introduced for SVMs



## Swiss roll Example





*Squashing by projecting onto a plane (left) versus unrolling the Swiss roll (right)*

# Kernel PCA: The Syntax

**Import the class containing the dimensionality reduction method**

```
from sklearn.decomposition import KernelPCA
```

**Create an instance of the class**

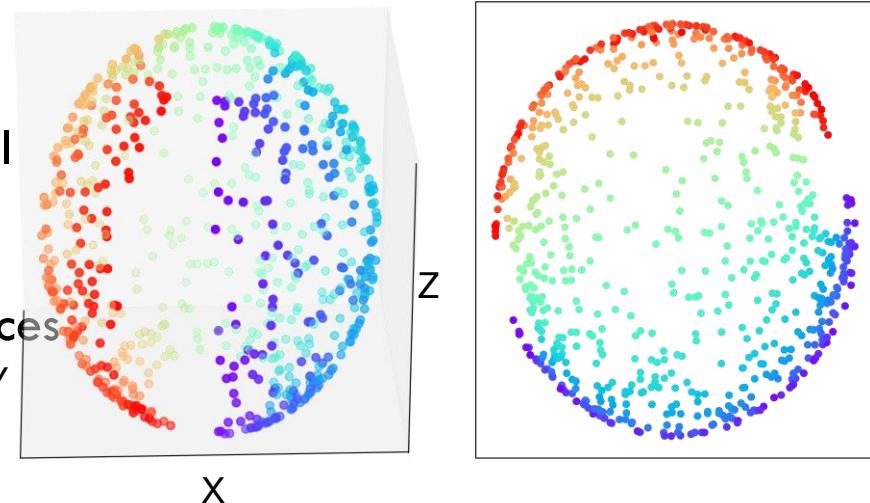
```
kPCA = KernelPCA(n_components=3, kernel='rbf', gamma=1.0)
```

**Fit the instance on the data and then transform the data**

```
X_trans = kPCA.fit_transform(X_train)
```

# Multi-Dimensional Scaling (MDS)

- Non-linear transformation
- Doesn't focus on maintaining overall variance
- Instead, maintains geometric distances between points



# MDS: The Syntax

**Import the class containing the dimensionality reduction method**

```
from sklearn.manifold import MDS
```

**Create an instance of the class**

```
mdsMod = MDS(n_components=2)
```

**Fit the instance on the data and then transform the data**

```
X_trans = mdsMod.fit_transform(X_sparse)
```

**Many other manifold dimensionality methods exist: Isomap, TSNE.**

# Uses of Dimensionality Reduction

- Frequently used for high dimensionality data
- Natural language processing (NLP)—many word combinations
- Image-based data sets—pixels are features



Image Source: [https://commons.wikimedia.org/wiki/File:Monarch\\_In\\_May.jpg](https://commons.wikimedia.org/wiki/File:Monarch_In_May.jpg)

# Uses of Dimensionality Reduction

- Divide image into  $12 \times 12$  pixel sections

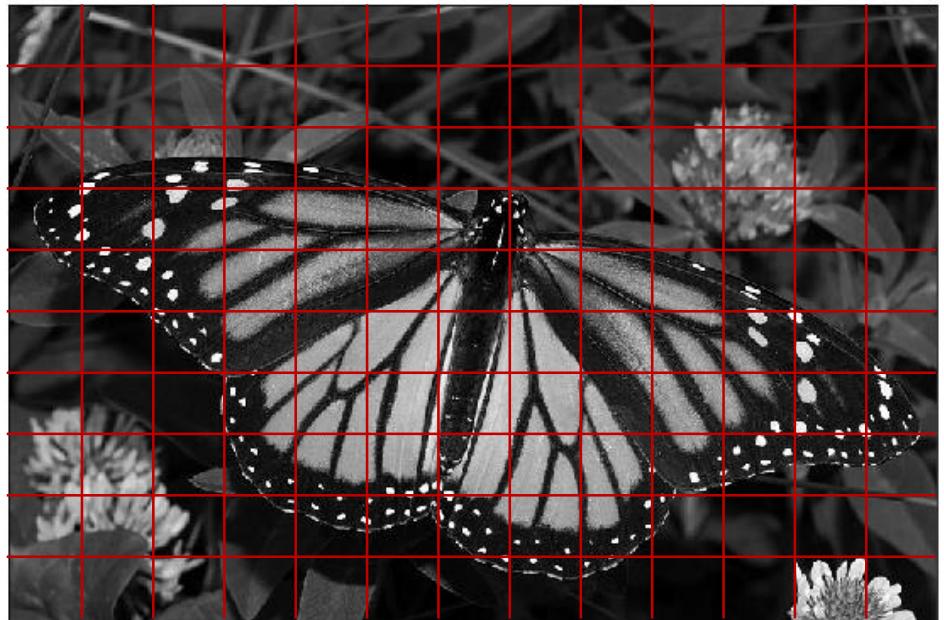


Image Source: [https://commons.wikimedia.org/wiki/File:Monarch\\_In\\_May.jpg](https://commons.wikimedia.org/wiki/File:Monarch_In_May.jpg)

# Uses of Dimensionality Reduction

- Divide image into  $12 \times 12$  pixel sections
- Flatten section to create row of data with 144 features

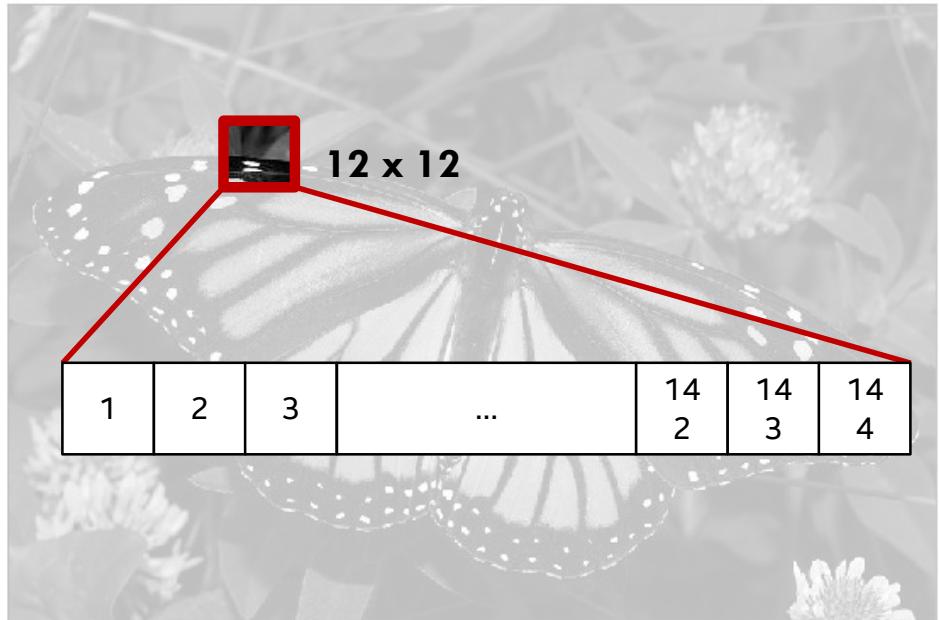
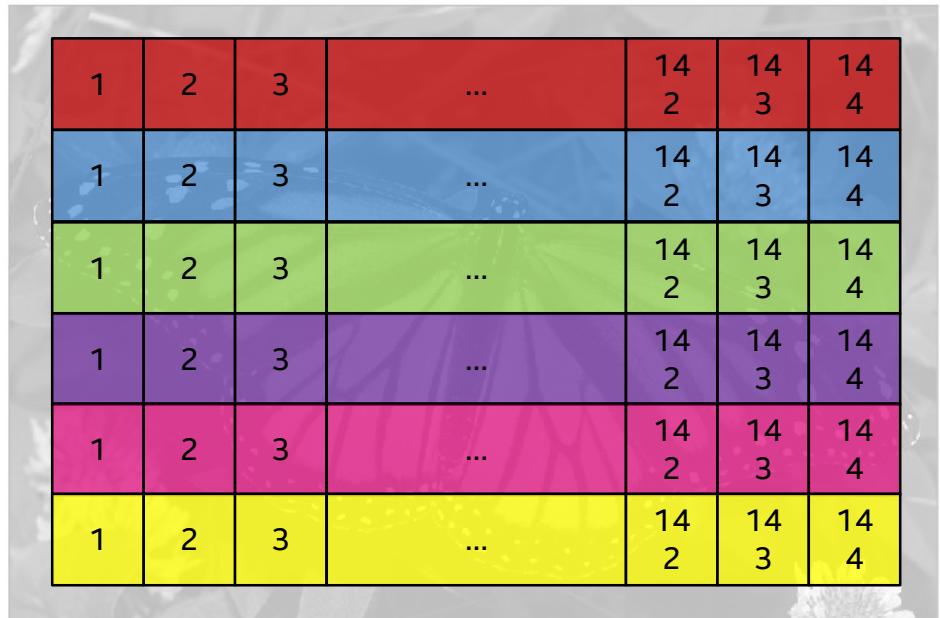


Image Source: [https://commons.wikimedia.org/wiki/File:Monarch\\_In\\_May.jpg](https://commons.wikimedia.org/wiki/File:Monarch_In_May.jpg)

# Uses of Dimensionality Reduction

- Divide image into  $12 \times 12$  pixel sections
- Flatten section to create row of data with 144 features
- Perform PCA on all data points



1	2	3	...	14 2	14 3	14 4
1	2	3	...	14 2	14 3	14 4
1	2	3	...	14 2	14 3	14 4
1	2	3	...	14 2	14 3	14 4
1	2	3	...	14 2	14 3	14 4
1	2	3	...	14 2	14 3	14 4

Image Source: [https://commons.wikimedia.org/wiki/File:Monarch\\_In\\_May.jpg](https://commons.wikimedia.org/wiki/File:Monarch_In_May.jpg)

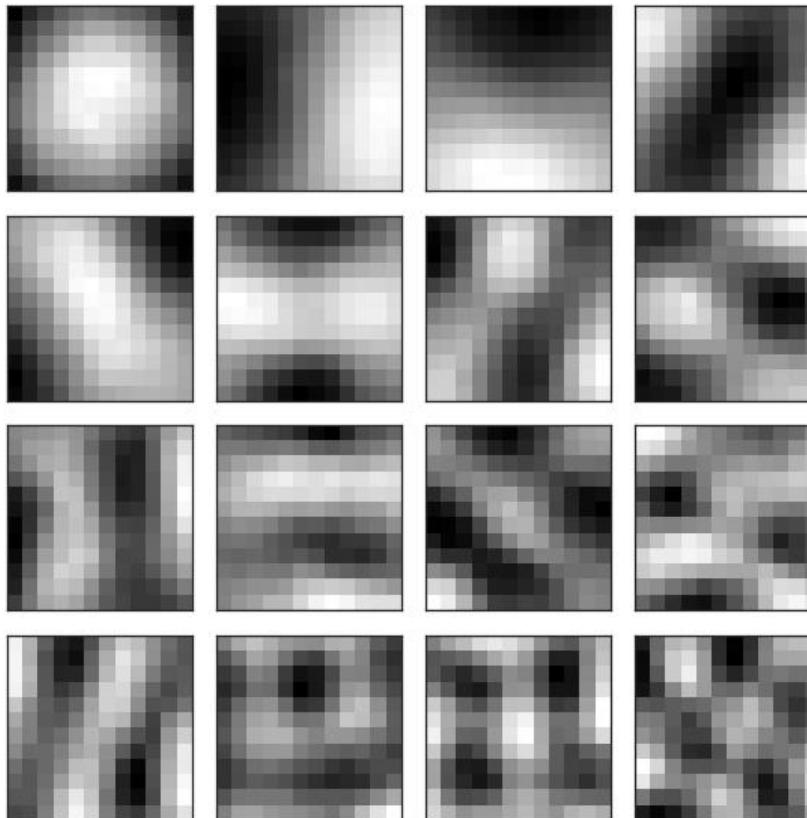
# PCA Compression: 144 → 60 Dimensions



# PCA Compression: 144 → 16 Dimensions



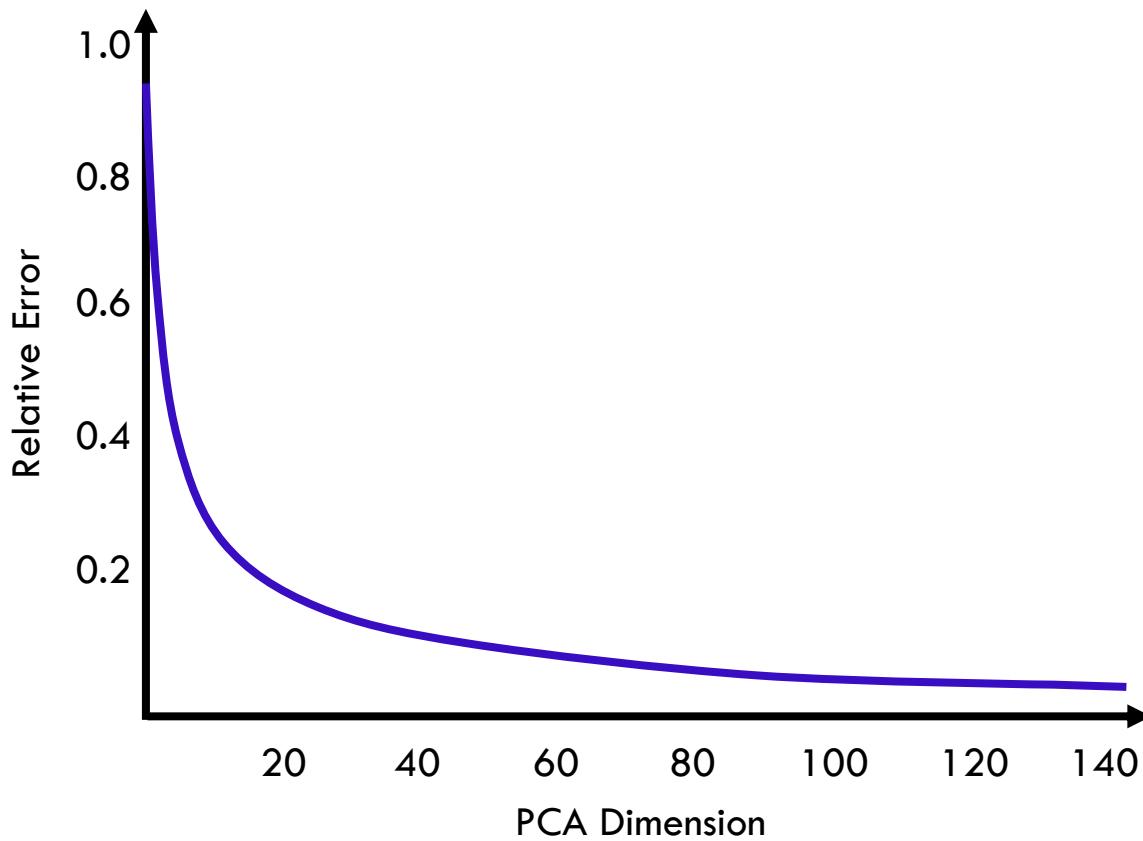
# Sixteen Most Important Eigenvectors



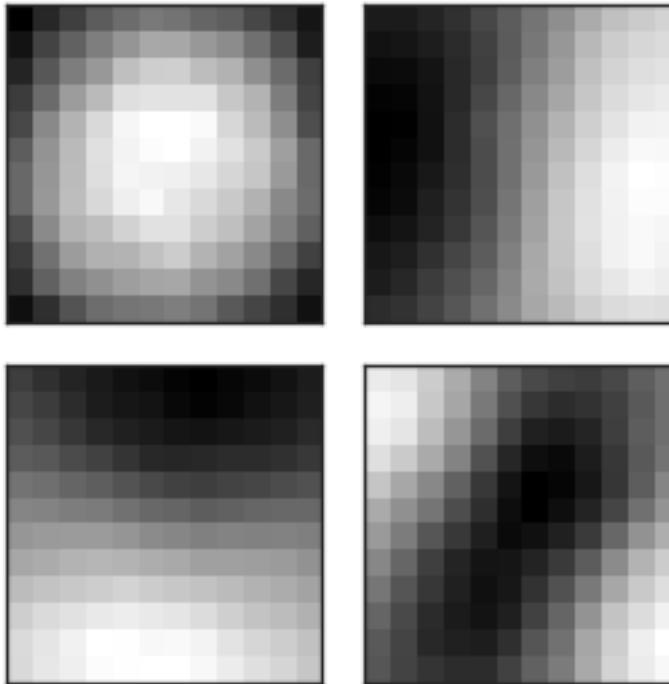
# PCA Compression: 144 → 4 Dimensions



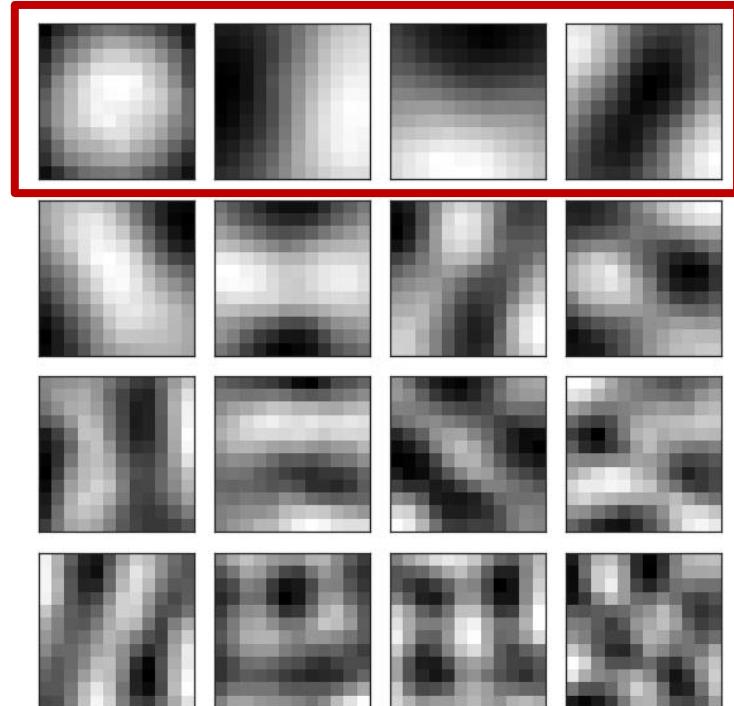
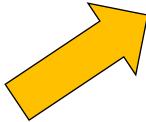
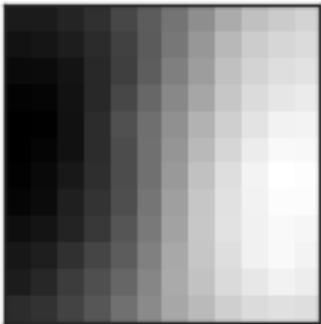
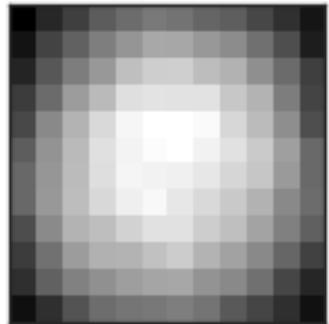
# L2 Error and PCA Dimension



# Four Most Important Eigenvectors



# Four Most Important Eigenvectors



# PCA Compression: 144 → 1 Dimension

