
ConvNets and Transfer Learning

Review

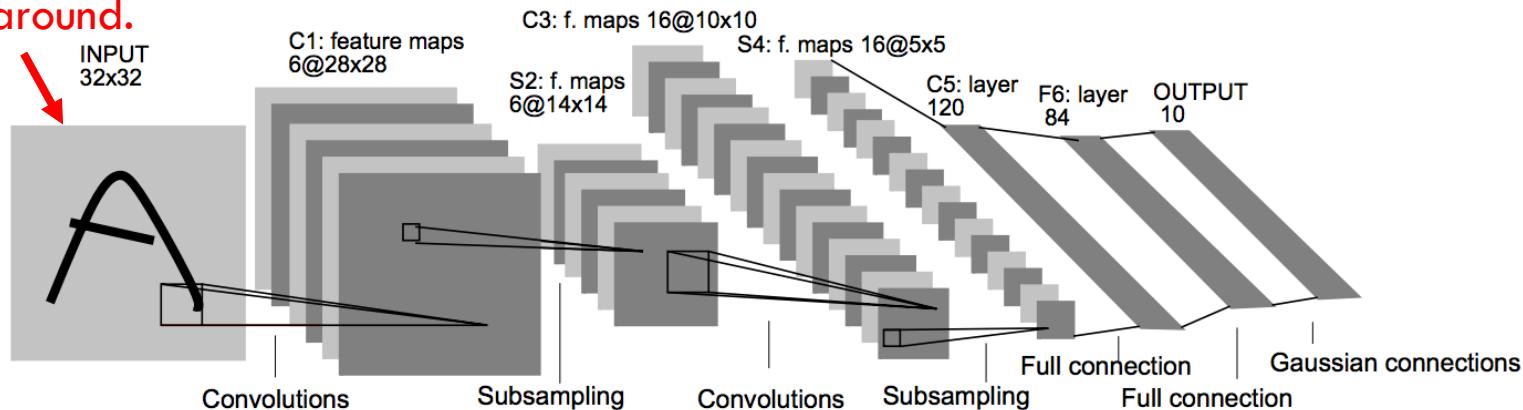
- Do some review of concepts from the last lecture
- We will revisit kernel, stride, and pooling in the context of the LeNet 5 model.

LeNet-5

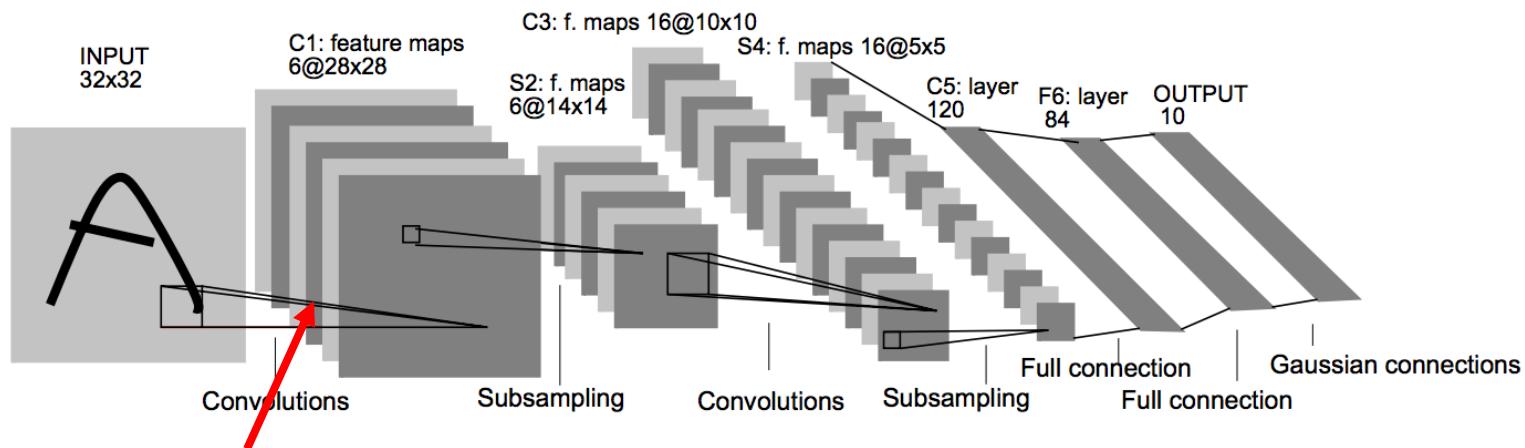
- Created by Yann LeCun in the 1990s
- Used on the MNIST data set.
- Novel Idea: Use convolutions to efficiently learn features on data set.

LeNet – Structure Diagram

Input: A 32×32 grayscale
image (28×28) with 2
pixels of padding all
around.

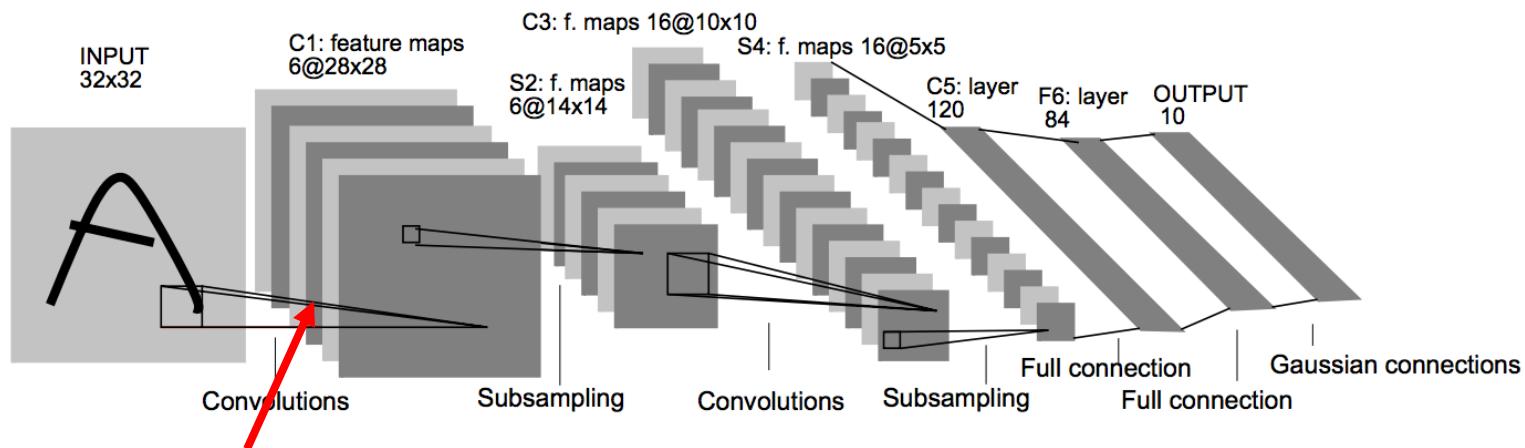


LeNet – Structure Diagram



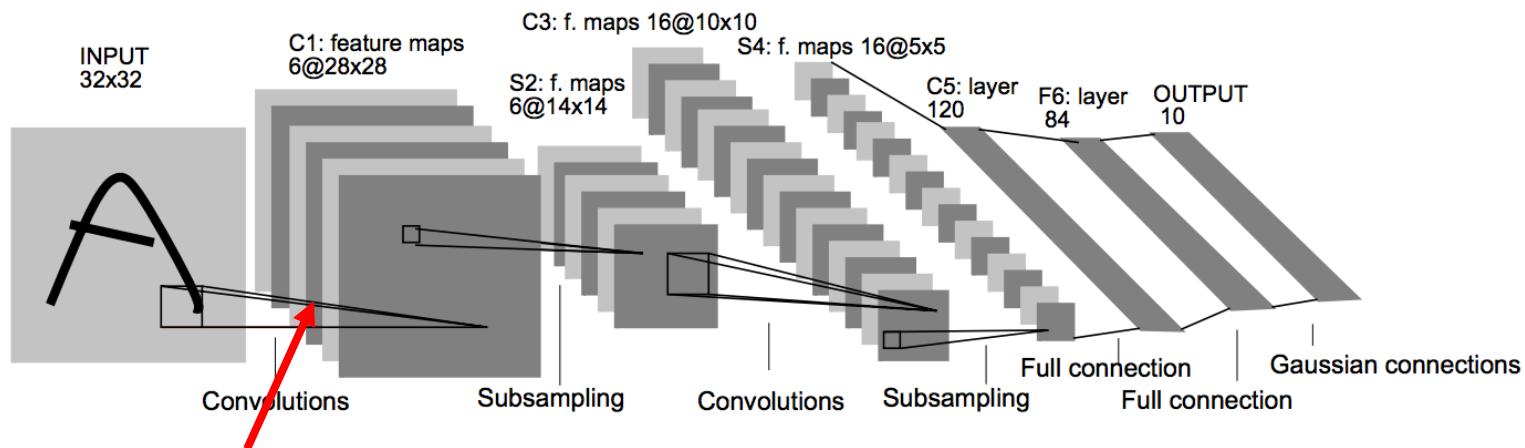
Next, we have a convolutional layer.

LeNet – Structure Diagram



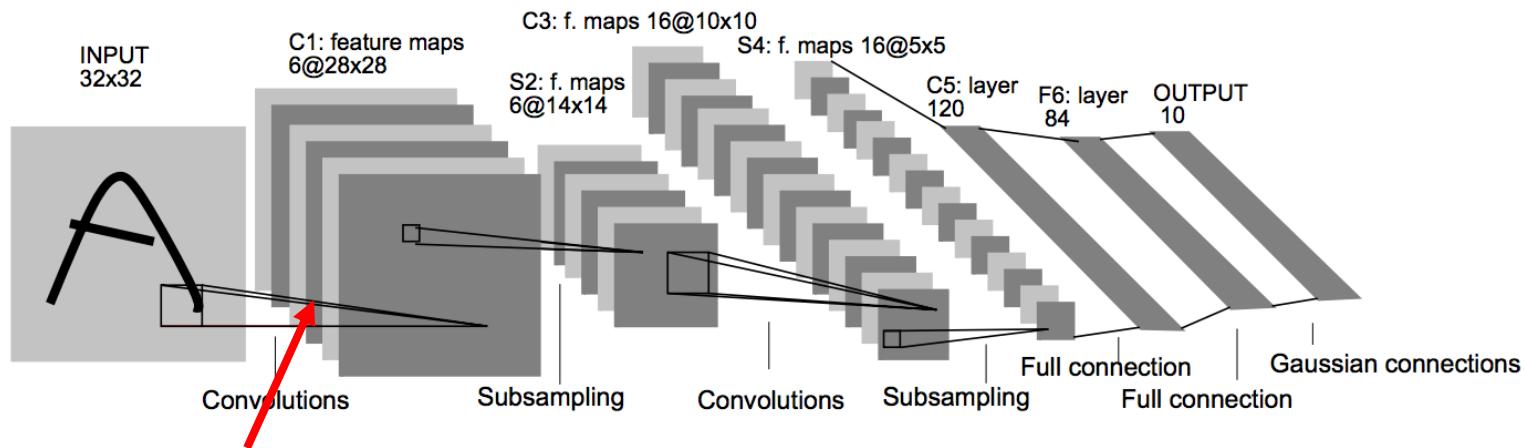
This is a **5x5 convolutional**
layer with stride 1.

LeNet – Structure Diagram



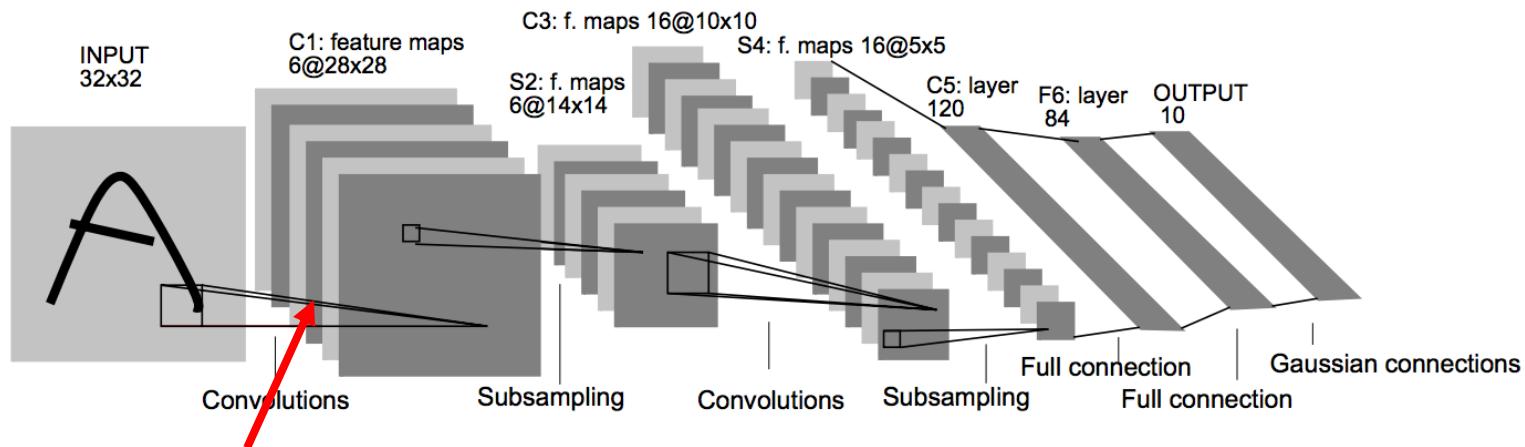
This means the resulting “filter” has dimension 28x28. (Why?)

LeNet – Structure Diagram



They use a depth of 6. This means
there are 6 different kernels that are
learned.

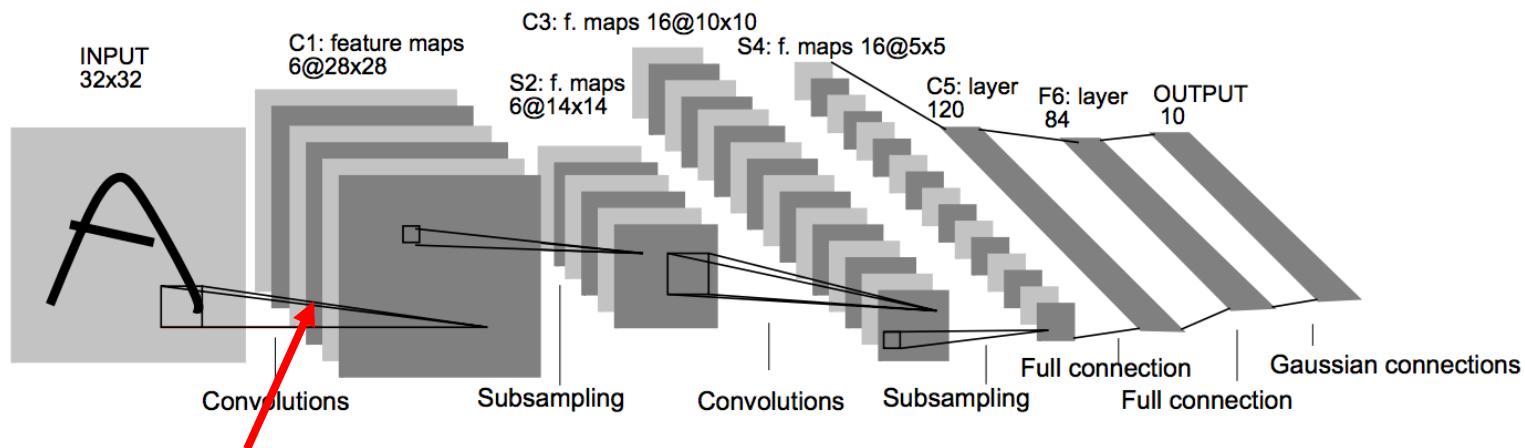
LeNet – Structure Diagram



They use a depth of 6. This means there are 6 different kernels that are learned.

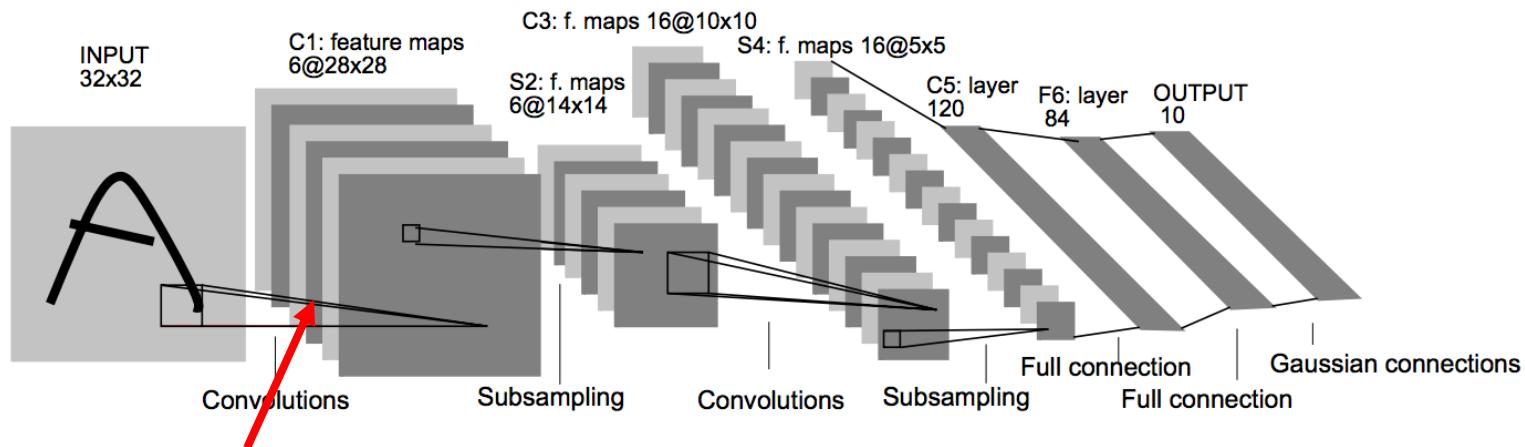
So the output of this layer is $6 \times 28 \times 28$.

LeNet – Structure Diagram



What is the total number of weights in this layer?

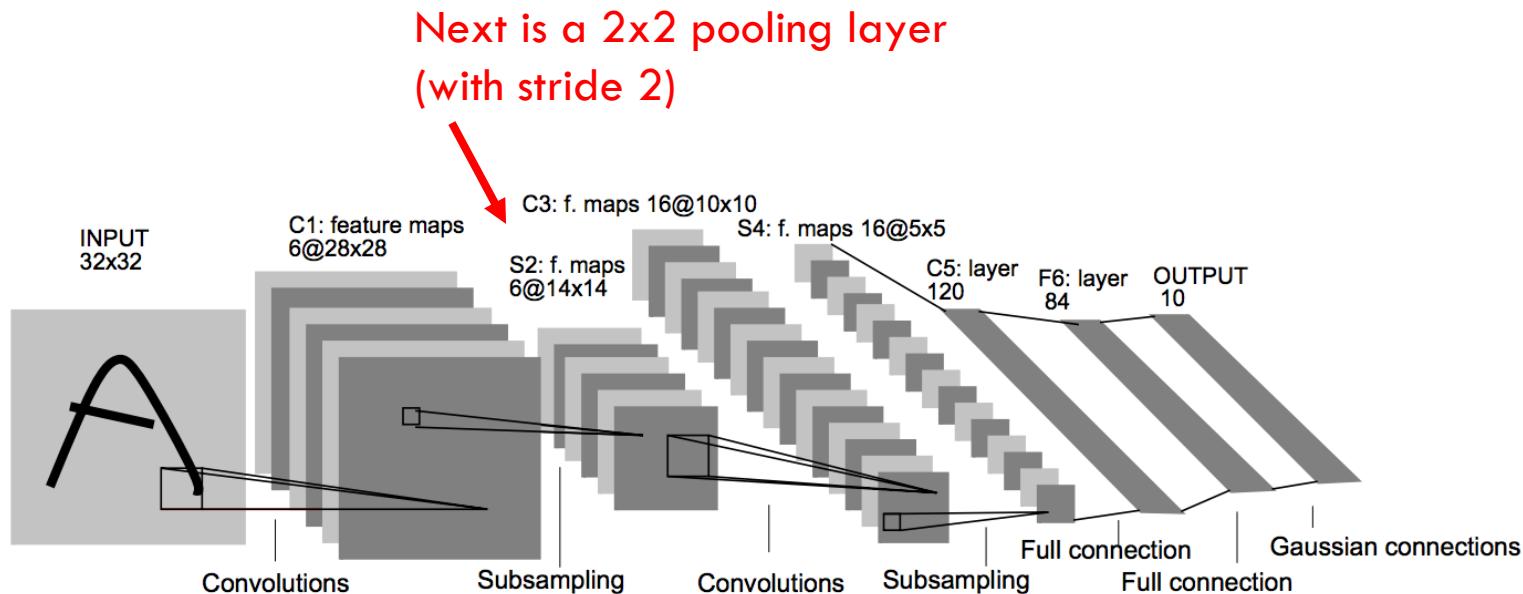
LeNet – Structure Diagram



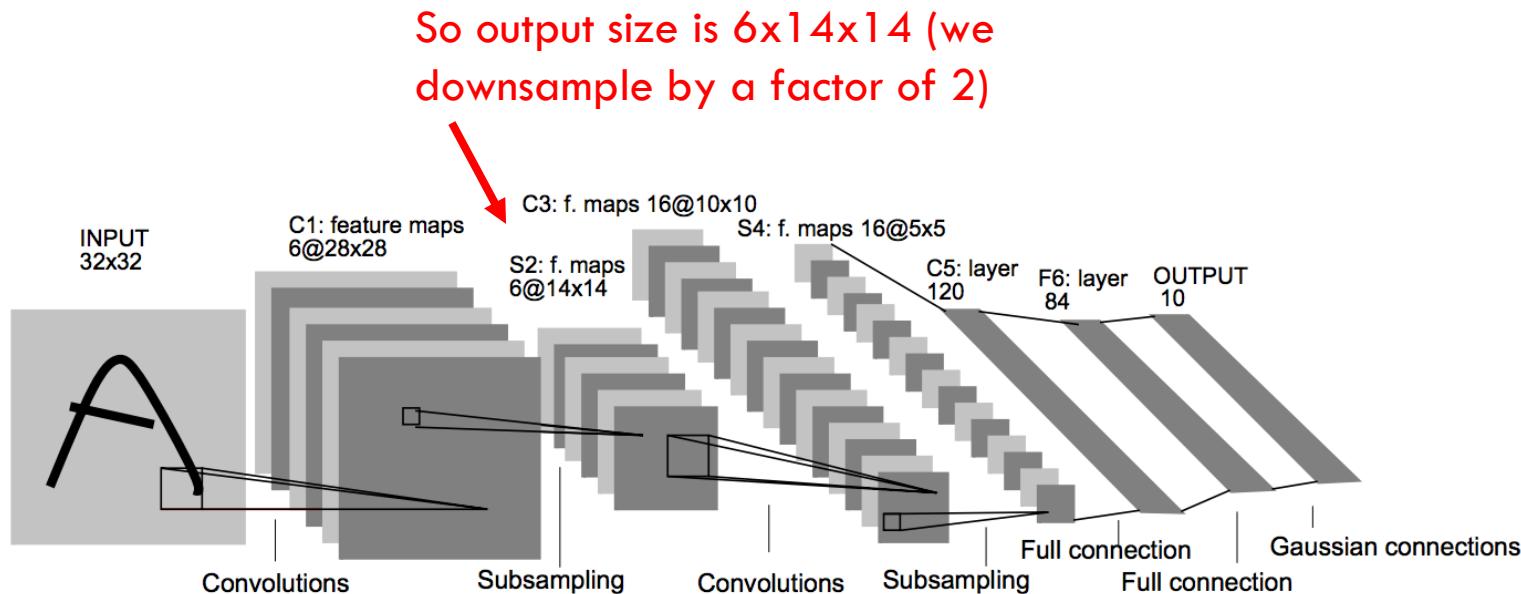
What is the total number of weights in this layer?

Answer: Each kernel has $5 \times 5 = 25$ weights (plus a bias term, so actually 26 weights). So total weights = $6 \times 26 = 156$.

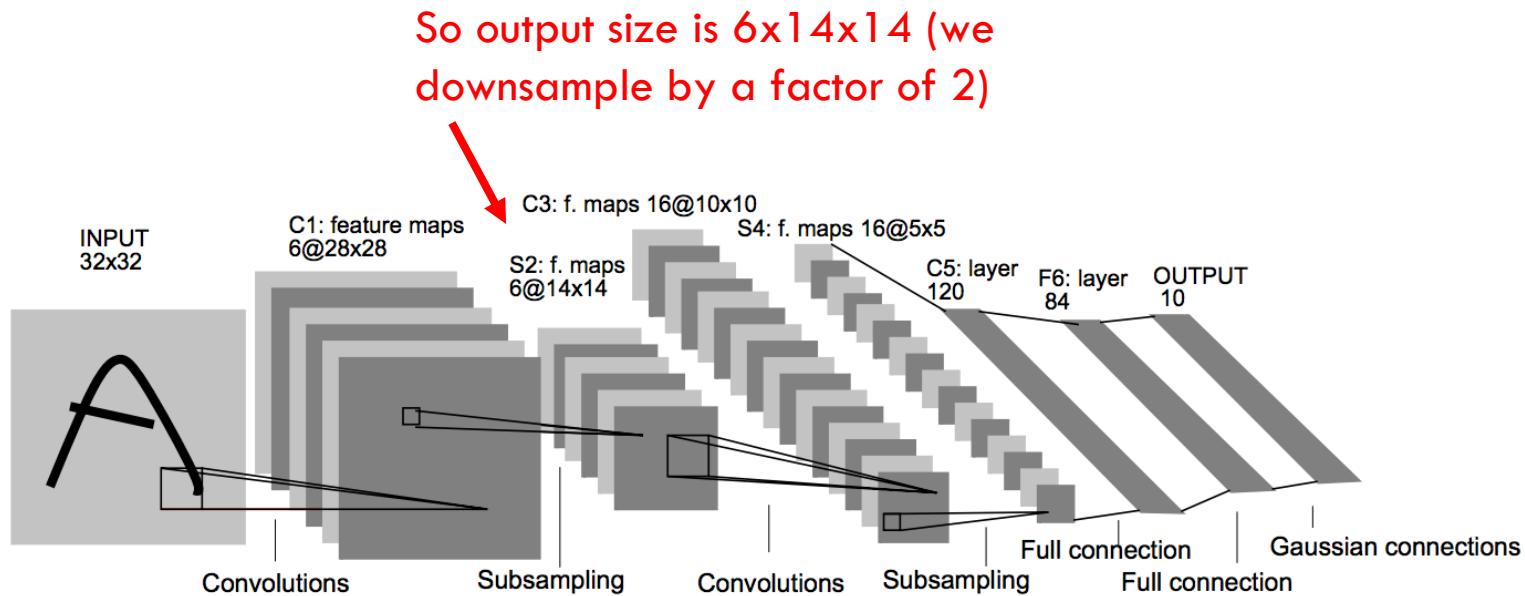
LeNet – Structure Diagram



LeNet – Structure Diagram



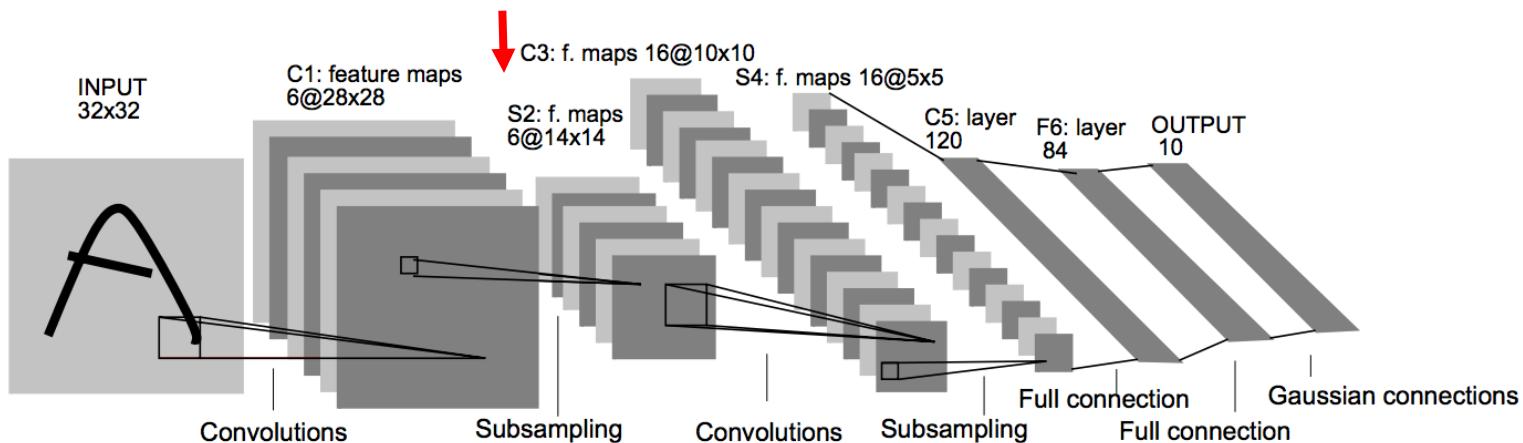
LeNet – Structure Diagram



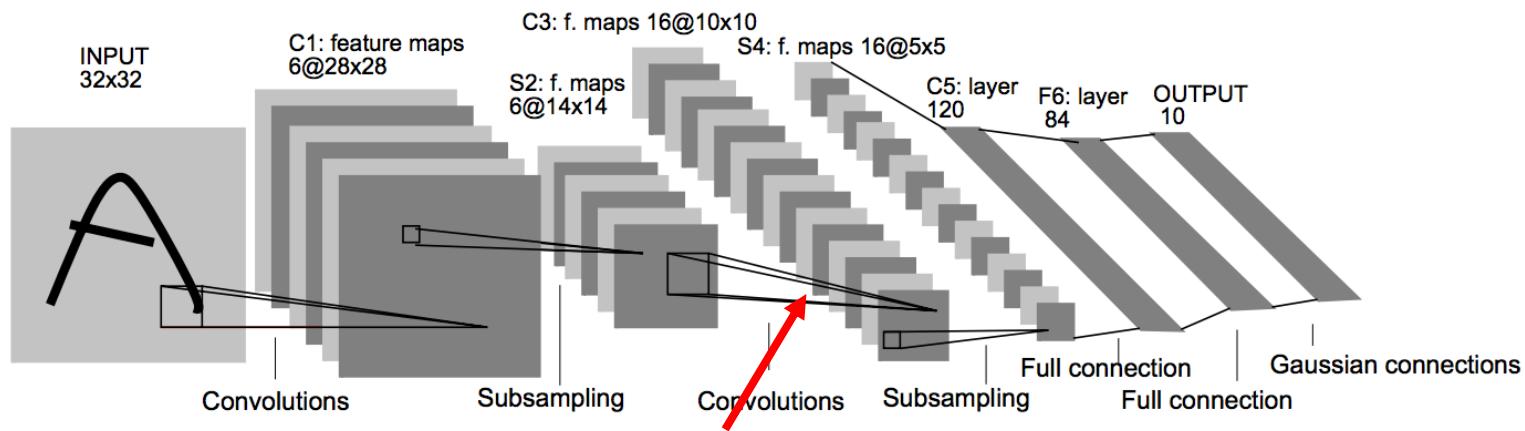
Note: The original paper actually does a more complicated pooling then max or avg pooling, but this is considered obsolete now.

LeNet – Structure Diagram

No weights! (pooling layers have no weights to be learned – it is a fixed operation.)

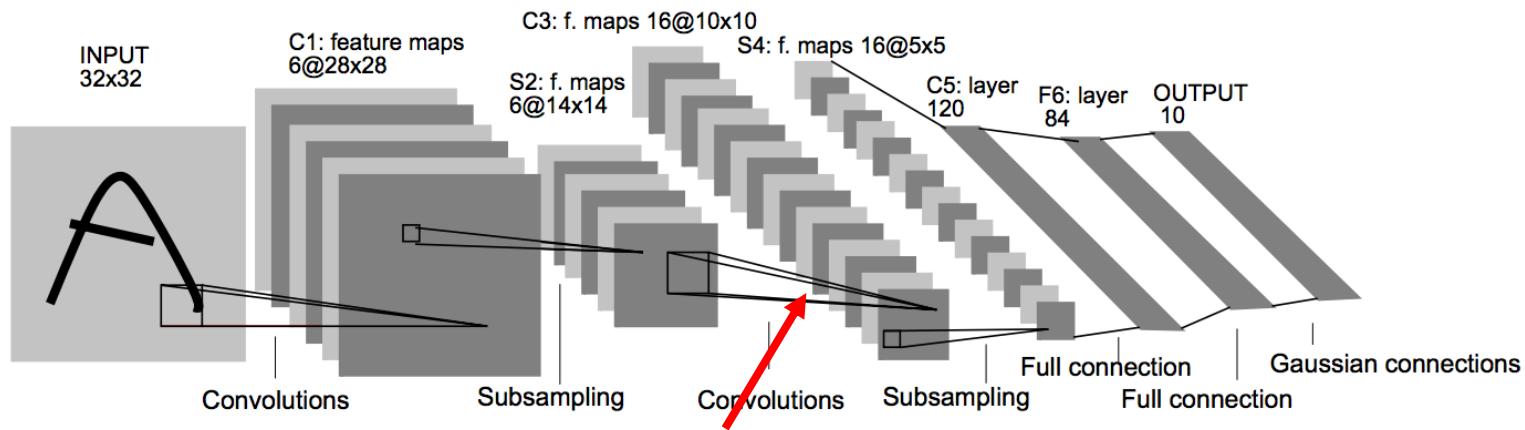


LeNet – Structure Diagram



Another 5x5 convolutional layer with stride 2.
This time the depth is 16.

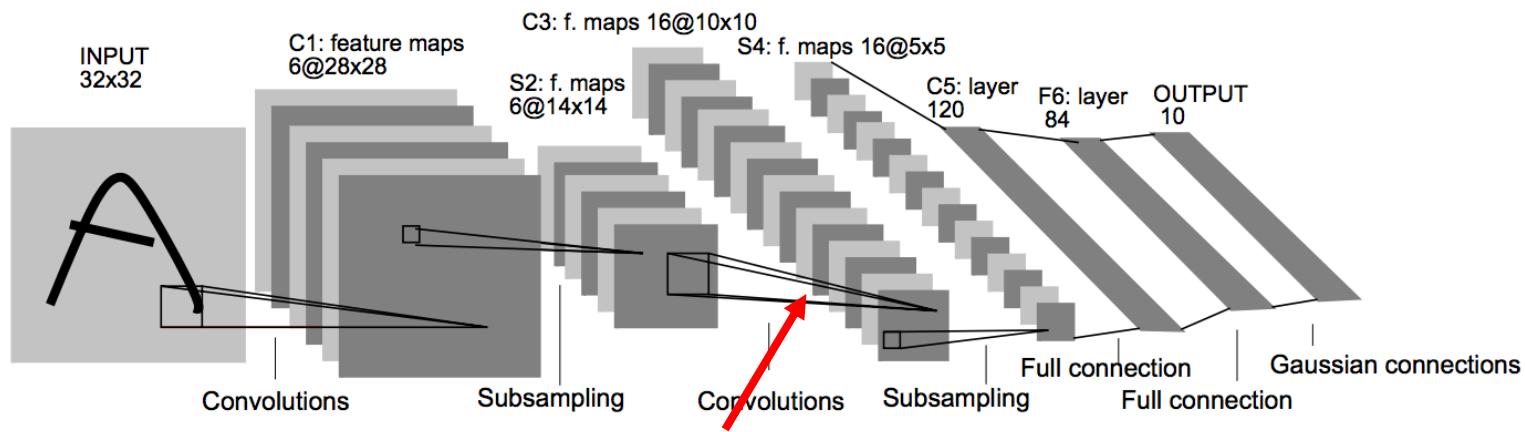
LeNet – Structure Diagram



Output size: $16 \times 10 \times 10$

How many weights? (tricky!)

LeNet – Structure Diagram

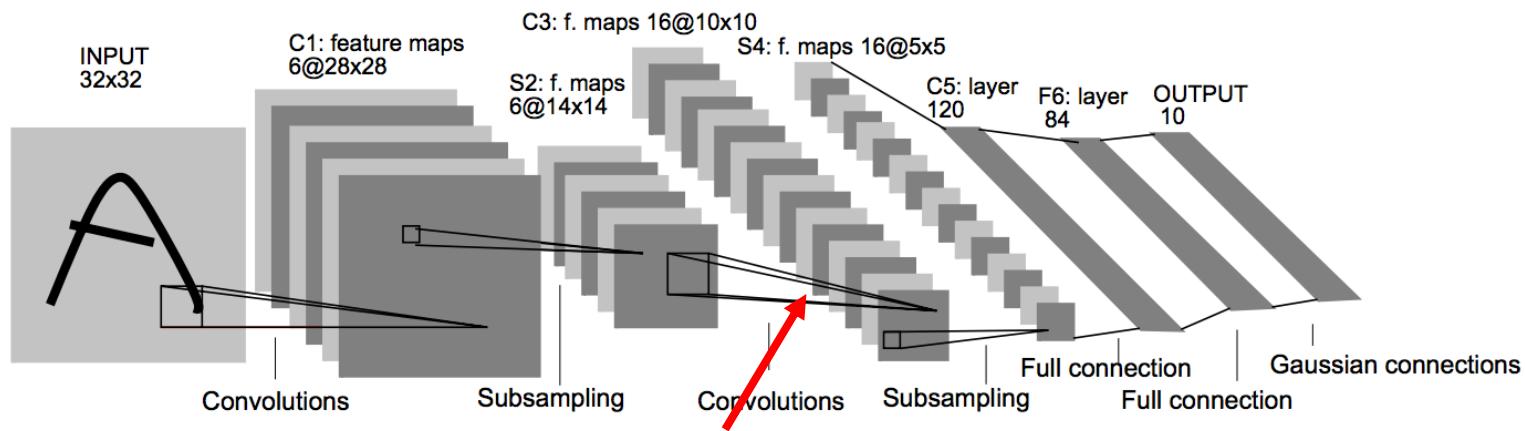


The kernels “take in” the full depth of the previous layer.

So each 5x5 kernel now “looks at” 6x5x5 pixels.

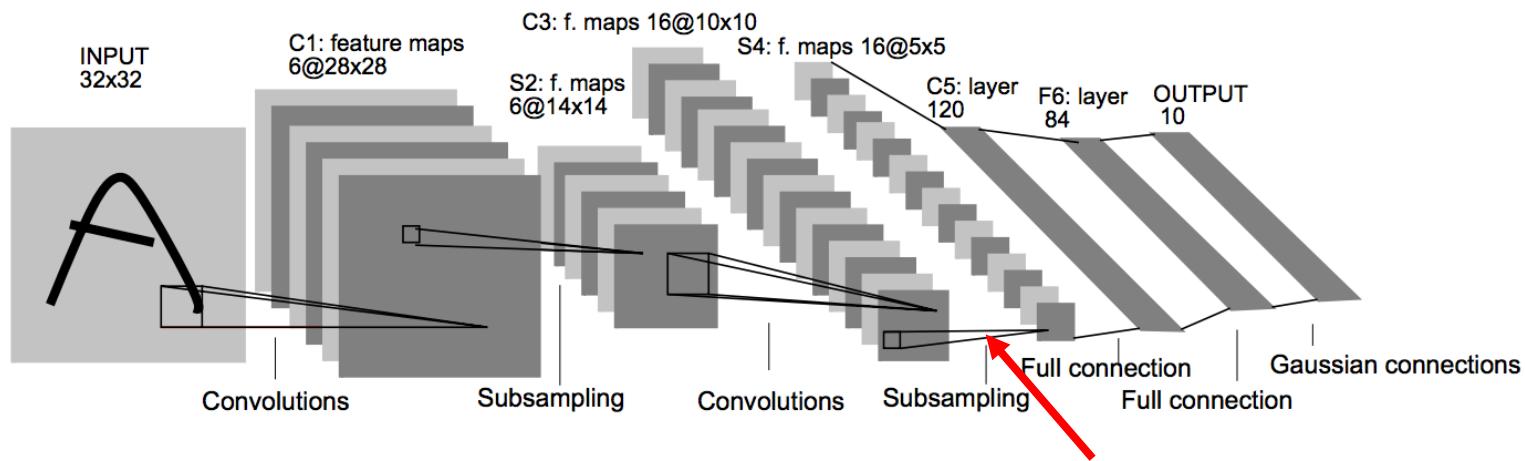
Each kernel has $6 \times 5 \times 5 = 150$ weights + bias term = 151.

LeNet – Structure Diagram



So, total weights for this layer = $16 * 151 = 2416$.

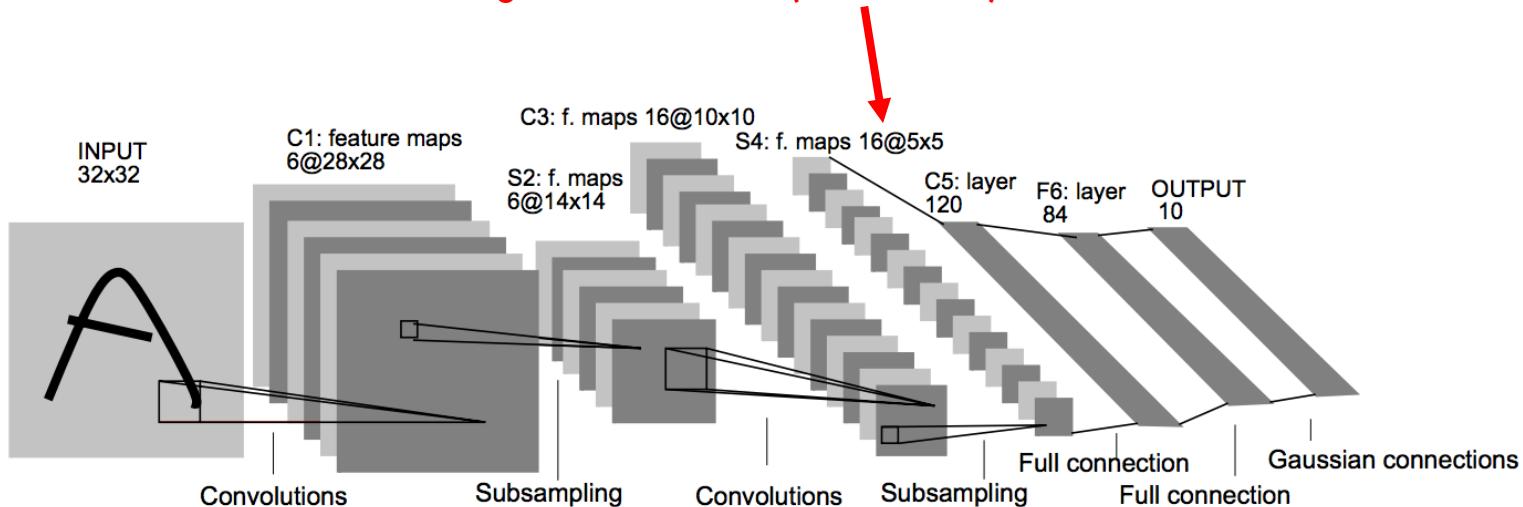
LeNet – Structure Diagram



Another 2x2 pooling layer.
Output is $16 \times 5 \times 5$.

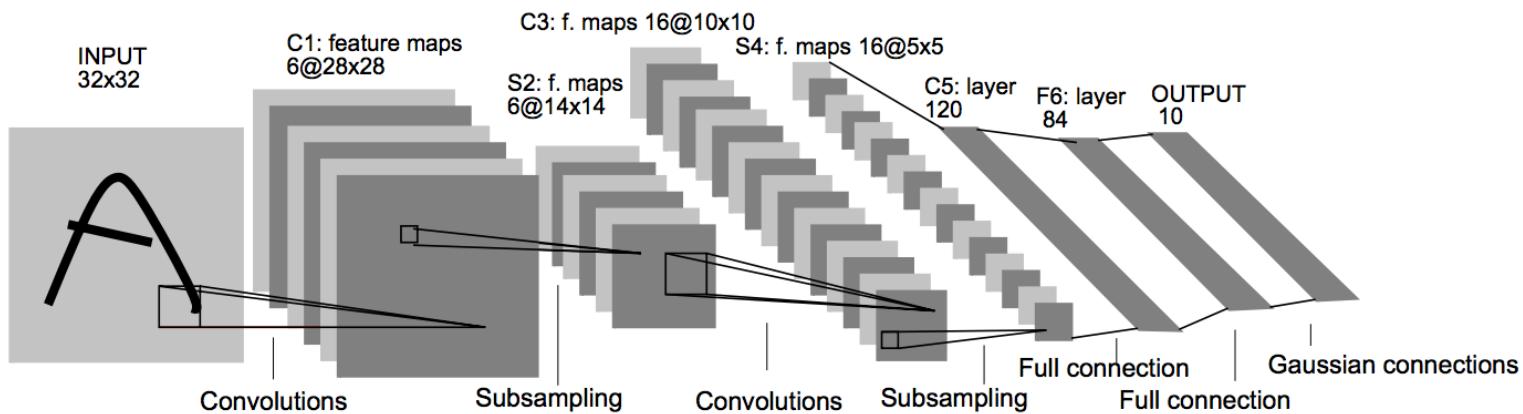
LeNet – Structure Diagram

We “flatten” this to a
length 400 vector (not shown)



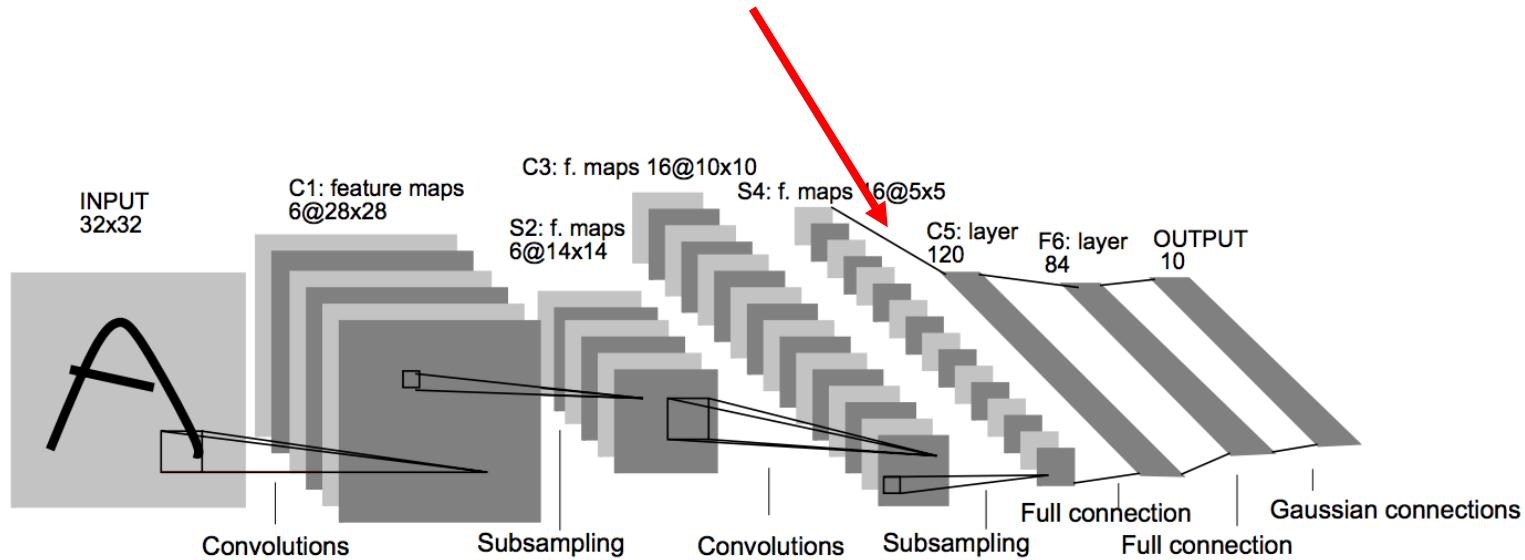
LeNet – Structure Diagram

The following layers are just
fully connected layers!



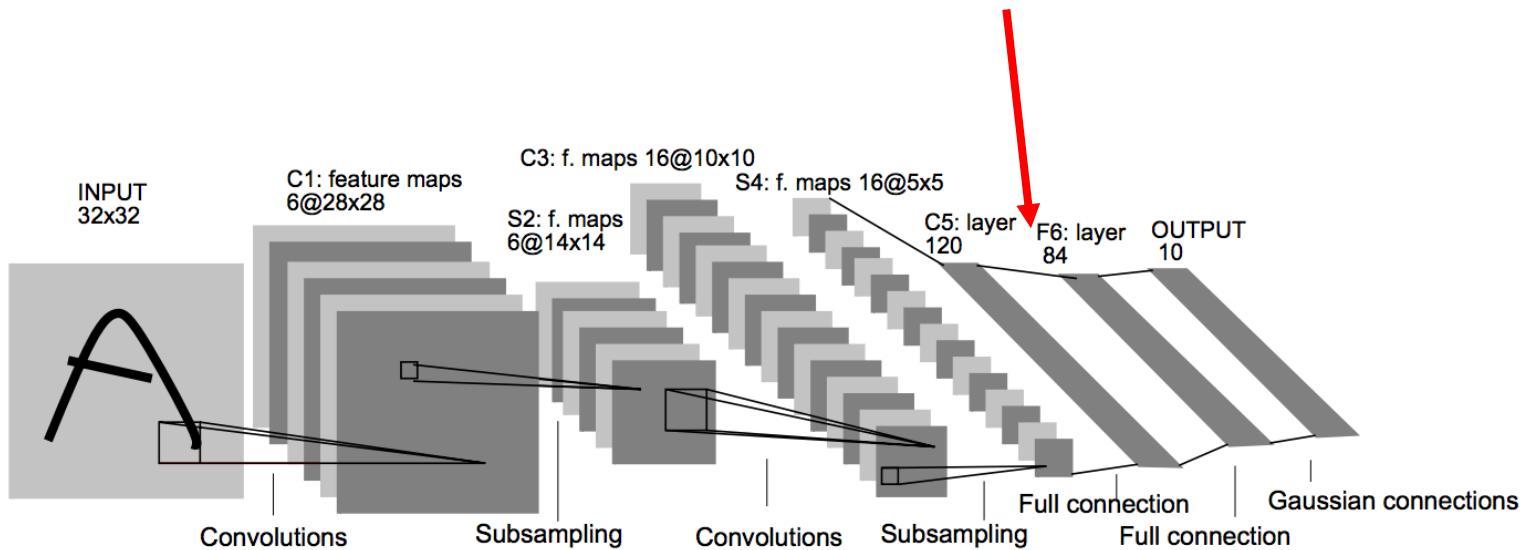
LeNet – Structure Diagram

From 400 to 120



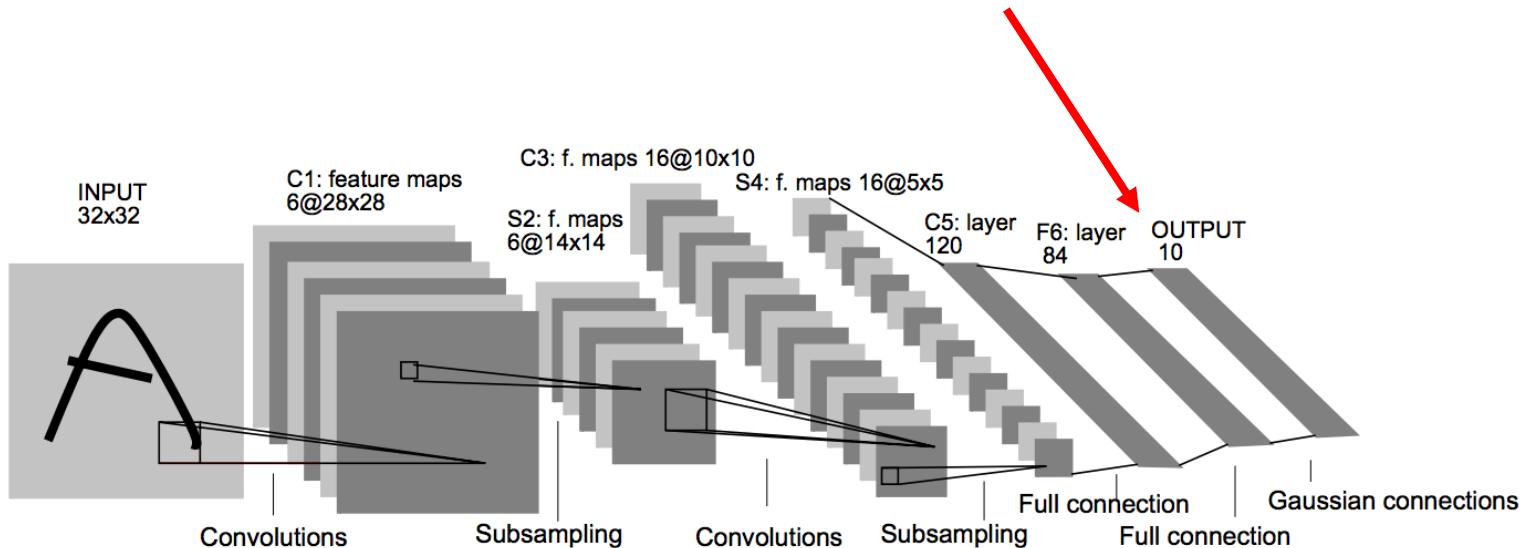
LeNet – Structure Diagram

Then from 120 to 84



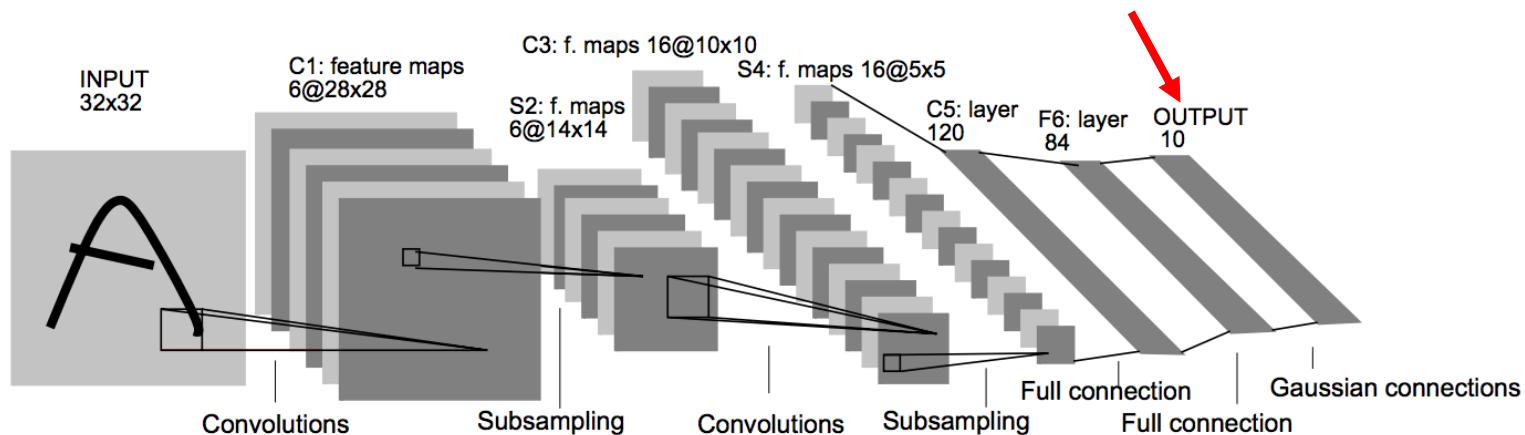
LeNet – Structure Diagram

Then from 84 to 10



LeNet – Structure Diagram

And a softmax output
of size 10 for the 10
digits



LeNet-5

How many total weights in the network?

$$\text{Conv1: } 1 \times 6 \times 5 \times 5 + 6 = 156$$

$$\text{Conv3: } 6 \times 16 \times 5 \times 5 + 16 = 2416$$

$$\text{FC1: } 400 \times 120 + 120 = 48120$$

$$\text{FC2: } 120 \times 84 + 84 = 10164$$

$$\text{FC3: } 84 \times 10 + 10 = 850$$

$$\text{Total: } = 61706$$

Less than a single FC layer with [1200x1200] weights!

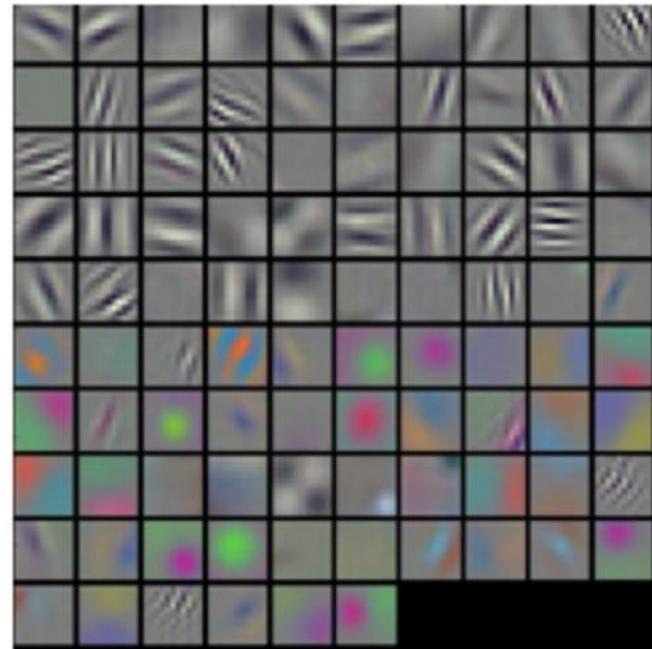
Note that Convolutional Layers have relatively few weights.

Transfer Learning



Motivation

- Early layers in a Neural Network are the hardest (i.e. slowest) to train
- Due to vanishing gradient property
- But these "primitive" features should be general across many image classification tasks



Motivation

- Later layers in the network are capturing features that are more particular to the specific image classification problem.
- Later layers are easier (quicker) to train since adjusting their weights has a more immediate impact on the final result.

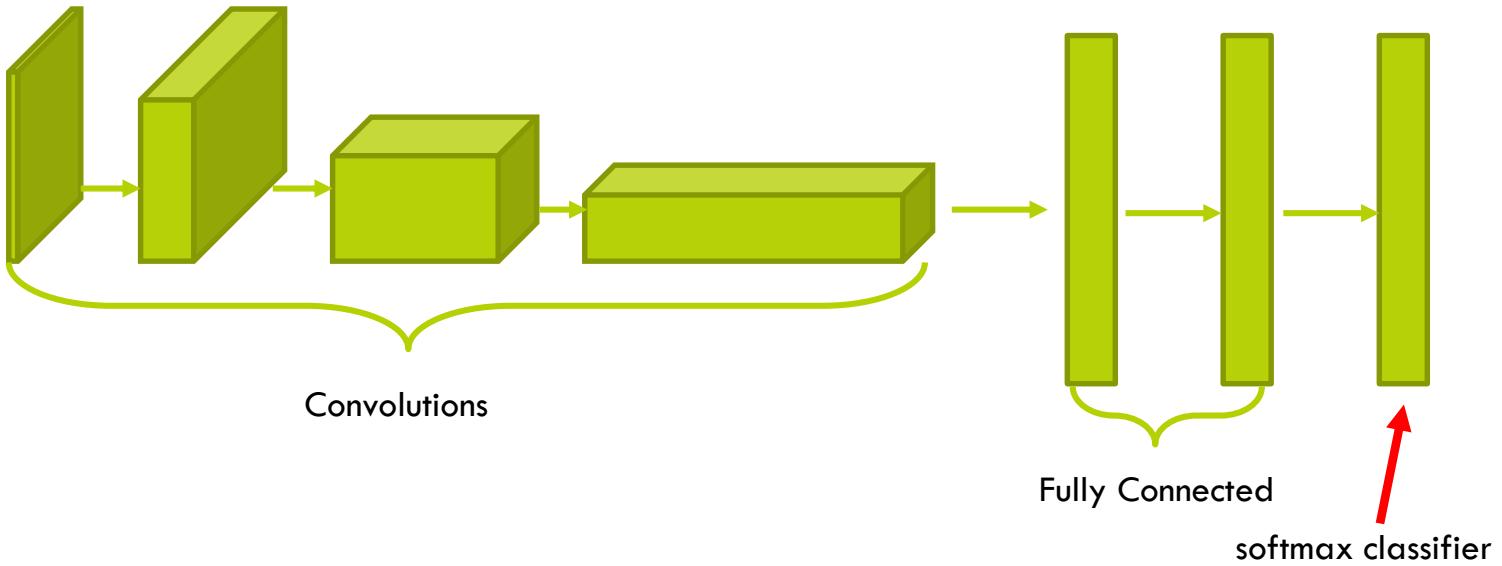
Motivation

- Famous, Competition-Winning Models are difficult to train from scratch
 - Huge datasets (like ImageNet)
 - Long number of training iterations
 - Very heavy computing machinery
 - Time experimenting to get hyper-parameters just right

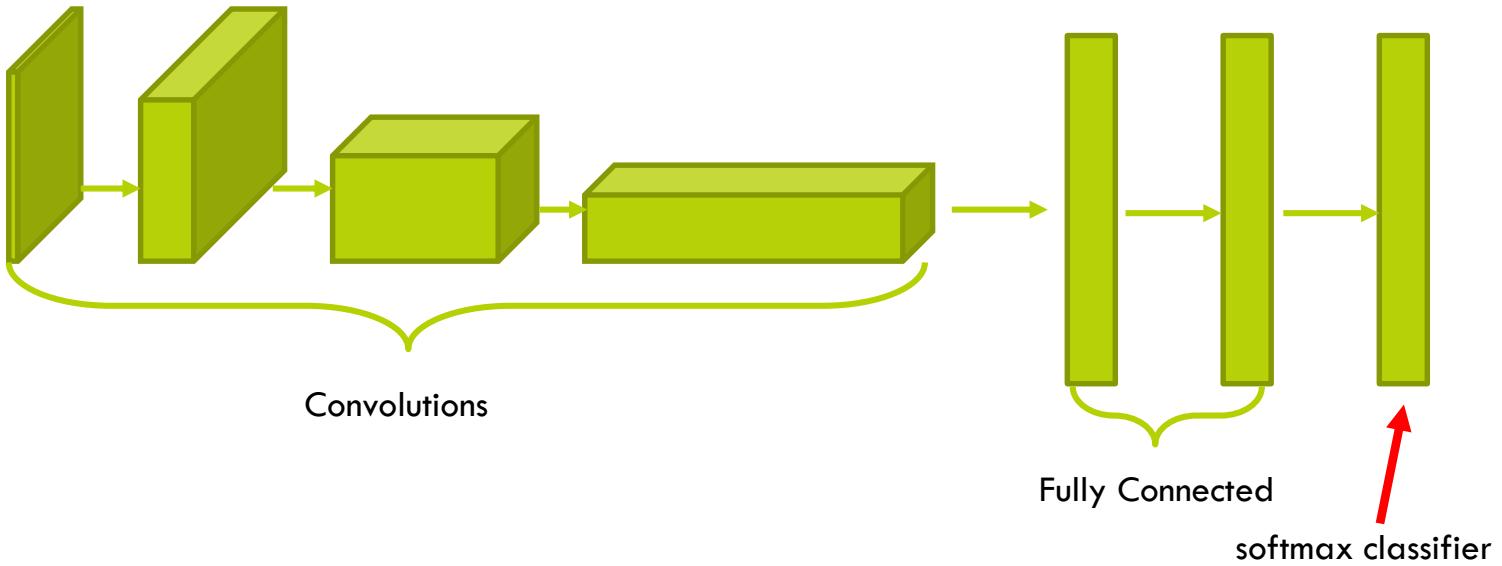
Transfer Learning

- However, the basic features (edges, shapes) learned in the early layers of the network *should* generalize.
- Results of the training are just weights (numbers) that are easy to store.
- Idea: keep the early layers of a pre-trained network, and re-train the later layers for a specific application
- This is called **Transfer Learning**.

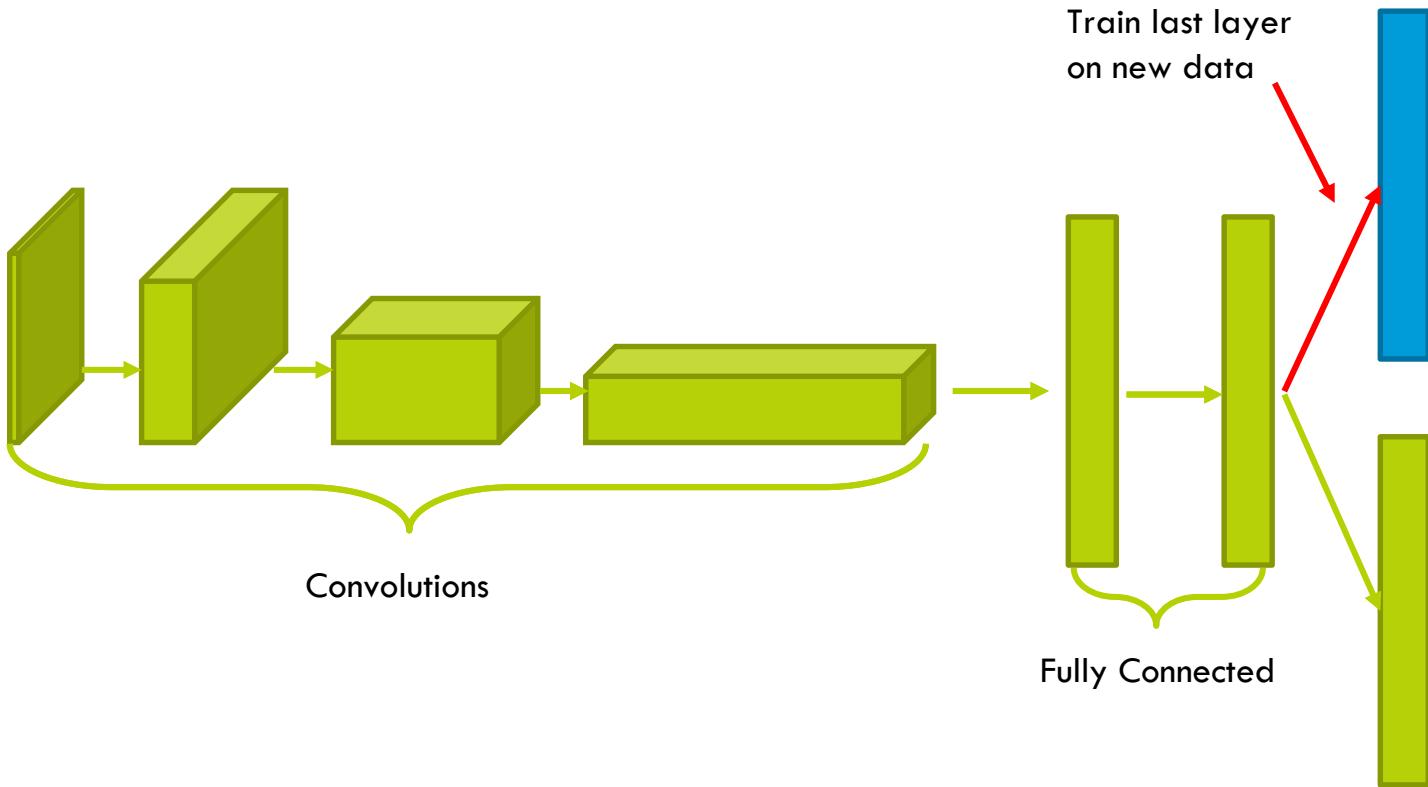
Transfer Learning



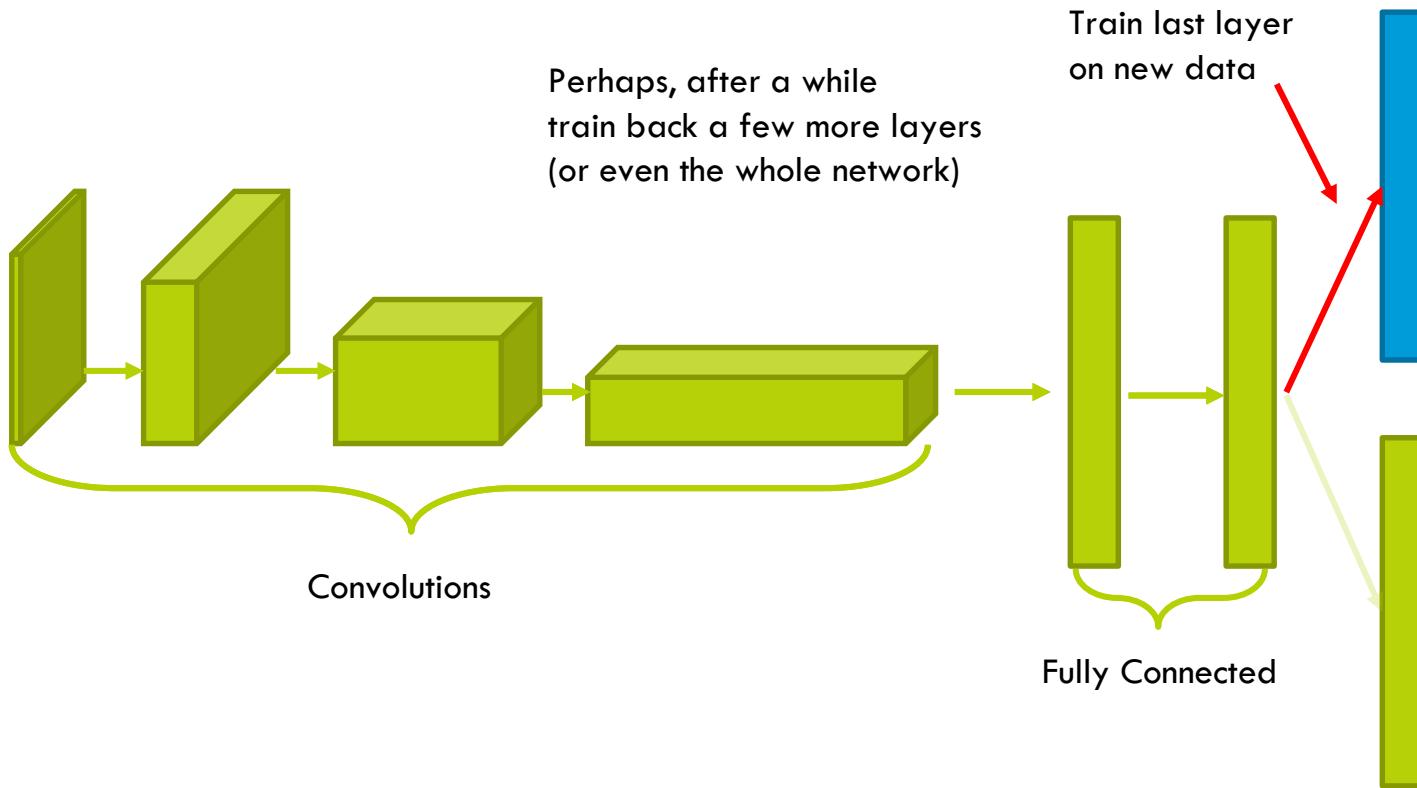
Transfer Learning



Transfer Learning



Transfer Learning



Transfer Learning Options

- The additional training of a pre-trained network on a specific new dataset is referred to as “Fine-Tuning”
- There are different options on “how much” and “how far back” to fine-tune.
 - Should I train just the very last layer?
 - Go back a few layers?
 - Re-train the entire network (from the starting point of the existing network)?

Guiding Principles for Fine-Tuning

- While there are no “hard and fast” rules, there are some guiding principles to keep in mind.
 - 1) The more similar your data and problem are to the source data of the pre-trained network, the less fine-tuning is necessary.
 - E.g. Using a network trained on ImageNet to distinguish “dogs” from “cats” should need relatively little fine-tuning. It already distinguished different breeds of dogs and cats, so likely has all the features you will need.

Guiding Principles for Fine-Tuning

- 2) The more data you have about your specific problem, the more the network will benefit from longer and deeper fine-tuning.
- E.g. If you have only 100 dogs and 100 cats in your training data, you probably want to do very little fine-tuning. If you have 10,000 dogs and 10,000 cats you may get more value from longer and deeper fine-tuning.

Guiding Principles for Fine-Tuning

- 3) If your data is substantially different in nature than the data the source model was trained on, Transfer Learning may be of little value.
- E.g. A network that was trained on recognizing typed Latin alphabet characters would not be useful in distinguishing cats from dogs. But it likely would be useful as a starting point for recognizing Cyrillic Alphabet characters.