

Convolutional Neural Net Architectures

AlexNet

- Created in 2012 for the ImageNet Large Scale Visual Recognition Challenge (ILSVRC)
- Task: predict the correct label from among 1000 classes
- Dataset: around 1.2 million images
- Considered the “flash point” for modern deep learning
- Demolished the competition.
- Top 5 error rate of 15.4%
- Next best: 26.2%

AlexNet - Model Diagram



AlexNet - Details

- They performed *data augmentation* for training
- Includes Cropping, horizontal flipping, and other manipulations

AlexNet - Details

- They performed *data augmentation* for training
 - Cropping, horizontal flipping, and other manipulations
- Basic Template:
 - Convolutions with ReLUs
 - Sometimes add maxpool after convolutional layer
 - Fully connected layers at the end before a softmax classifier


VGG

- Simplify Network Structure
- Avoid Manual Choices of Convolution Size
- Very Deep Network with 3x3 Convolutions
- These “effectively” give rise to larger convolutions


Reference:

Very Deep Convolutional Networks for Large-Scale Image Recognition
Karen Simonyan and Andrew Zisserman, 2014


VGG16 Diagram




VGG



VGG




VGG




VGG

We can say that the “receptive field” of Layer 2 is 3x3
Each output has been influenced by a 3x3 patch of inputs



VGG


What about on Layer 3?




VGG

This output on Layer 3 uses a 3x3 patch from Layer 2


How much from Layer 1 does it use?




VGG




VGG




VGG




VGG




VGG




VGG




VGG



VGG




VGG




VGG


Each square in Layer 3 “sees” a 5x5 grid from Layer 1




Layer 1
(Input)



Layer 2



Layer 3



VGG

Two 3x3, stride 1 convolutions in a row → one 5x5

Three 3x3 convolutions → one 7x7 convolution

Benefit: fewer parameters

One 3x3 layer

$$3 \times 3 \times C \times C = 9C^2$$

One 7x7 layer

$$7 \times 7 \times C \times C = 49C^2$$

Three 3x3 layers

$$3 \times (9C^2) = 27C^2$$

$$49C^2 \rightarrow 27C^2 \rightarrow \approx 45\% \text{ reduction!}$$

VGG


- One of the first architectures to experiment with many layers
(More is better!)
- Can use multiple 3x3 convolutions to simulate larger kernels with fewer parameters
- Served as "base model" for future works

Inception


- Szegedy et al 2014
- Idea: network would want to use different receptive fields
- Want computational efficiency
- Also want to have sparse activations of groups of neurons
- Hebbian principle: “Fire together, wire together”
- Solution: Turn each layer into branches of convolutions
- Each branch handles smaller portion of workload
- Concatenate different branches at the end

Inception


Basic idea: replace single 3×3 convolutions with module




Inception



Inception




Inception




Inception


We also control the number of filters from max pool branch with a 1x1 conv



Inception





Inception V3 schematic



ResNet - Motivation

Issue: Deeper Networks performing worse on **training** data!
(as well as test data)




ResNet

- Surprising because deeper networks should overfit more.
- So what's happening?
- Early layers of Deep Networks are very slow to adjust.
- Analogous to “Vanishing Gradient” issue.
- In theory, should be able to just have an “identity” transformation that makes the deeper network behave like a shallower one

ResNet

- Assumption: best transformation over multiple layers is close to $\mathcal{F}(x) + x$
- $x \rightarrow$ input to series of layers
- $\mathcal{F}(x) \rightarrow$ function represented by several layers (such as convs)
- Enforce this by adding “shortcut connections”
- Add the inputs from an earlier layer to the output of current layer



ResNet

- Add previous layer back in to current layer!
- Similar idea to “boosting”

