



Software

# Advanced Techniques for CNNs and Keras

# Data Augmentation

- One practical obstacle to building image classifiers is obtaining labeled training data.
- Finding images is difficult.
- Labeling images is time consuming and costly.
- How can we make the most out of the labelled data we have?

# Data Augmentation

- If this is a chair:



# Data Augmentation

- If this is a chair...



- Then so is this!



# Data Augmentation

- If this is a chair...



- Also this:



# Data Augmentation

- If this is a chair...



- Also this:



# Data Augmentation

- By slightly altering images, we can increase our effective data size.
- Also allows the neural network to learn invariance to certain transformations.
- But we need to be careful – this can have unintended consequences.

# Data Augmentation

Would not want a self-driving car to think these mean the same thing!



# Data Flows in Keras

- Keras has a convenient mechanism for Data Augmentation.
- Requires use of “Data Generators”
- To date, we have used the standard `model.fit` mechanism
- Holds entire dataset in memory
- Reads the batches one by one out of memory

# Data Flows in Keras

- Alternative mechanism is to use a “data generator”
- Idea: define a generator object which “serves” the batches of data.
- Then use `model.fit_generator` instead of `model.fit`
- Generators can be used to serve images from disk to conserve working memory

# ImageDataGenerator

- Keras has an `ImageDataGenerator` class which permits “real-time” data-augmentation.
- When a batch of images is served, you can specify random changes to be made to the image.
- These include shifting, rotating, flipping, and various normalizations of the pixel values.

# ImageDataGenerator

```
keras.preprocessing.image.ImageDataGenerator(  
    featurewise_center=False, samplewise_center=False,  
    featurewise_std_normalization=False, samplewise_std_normalization=False,  
    zca_whitening=False,  
    rotation_range=0.,  
    width_shift_range=0.,  
    height_shift_range=0.,  
    shear_range=0., zoom_range=0., channel_shift_range=0.,  
    fill_mode='nearest', cval=0.,  
    horizontal_flip=False, vertical_flip=False,  
    rescale=None, preprocessing_function=None,  
    data_format=K.image_data_format())
```

# Shifting Images

```
keras.preprocessing.image.ImageDataGenerator(  
    width_shift_range=0.,  
    height_shift_range=0.,  
    ...)
```

`width_shift_range=0.1`

# Shifting Images (how to fill in)

```
keras.preprocessing.image.ImageDataGenerator(  
...,  
fill_mode='nearest', cval=0.,  
...)
```

cval "constant", "nearest", "reflect", "wrap"

# Rotating Images

```
keras.preprocessing.image.ImageDataGenerator(  
...,  
rotation_range=0.,  
...)
```

`rotation_range=30`

# Flipping Images


```
keras.preprocessing.image.ImageDataGenerator(  
...,  
horizontal_flip=False, vertical_flip=False,  
...)
```

# Keras Functional API

- So far we have primarily used the Keras Sequential method.
- Very convenient when each additional layer takes the results of the previous layer.
- However, suppose we have more than one input to a particular layer.
- Suppose we have multiple outputs, and want different loss functions for these outputs.
- More complicated graph structures require the Functional API.

# Keras Functional API - principles

- Every layer is “callable” on a tensor, and returns a tensor.
- Specifically designate inputs using the Input layer.
- Merge outputs into a single output using  
`keras.layers.concatenate`
- Stack layers in a similar fashion to the Sequential model.
- Use `Model` to specify inputs and outputs of your model.
- When compiling, can specify different losses for different outputs, and specify how they should be weighted.



Software